

U-MATH: A UNIVERSITY-LEVEL BENCHMARK FOR EVALUATING MATHEMATICAL SKILLS IN LLMs

Konstantin Chernyshev*, Vitaliy Polshkov, Ekaterina Artemova, Sergei Tilga

Toloka AI

{kchernyshev, cogwheelhead, katya-art, tilgasergey}@toloka.ai

Alex Myasnikov, Vlad Stepanov

Gradarius

{alex, vstepanov}@gradarius.com

Alexei Miasnikov

Gradarius, Stevens Institute of Technology

amiasnik@stevens.edu

ABSTRACT

The current evaluation of mathematical skills in LLMs is limited, as existing benchmarks are either relatively small, primarily focus on elementary and high-school problems, or lack diversity in topics. Additionally, the inclusion of visual elements in tasks remains largely under-explored.

To address these gaps, we introduce **U-MATH**, a novel benchmark of **1,100** unpublished open-ended university-level problems sourced from teaching materials. It is balanced across six core subjects, with **20% of multimodal problems**. Given the open-ended nature of U-MATH problems, we employ an LLM to judge the correctness of generated solutions. To this end, we release **μ -MATH**, a dataset to evaluate the LLMs' capabilities in judging solutions.

The evaluation of general domain, math-specific, and multimodal LLMs highlights the challenges presented by U-MATH. Our findings reveal that LLMs achieve a maximum accuracy of only 63% on text-based tasks, with even lower 45% on visual problems. The solution assessment proves challenging for LLMs, with the best LLM judge having an F1-score of 80% on **μ -MATH**.

We open-source U-MATH, **μ -MATH**, and evaluation code on GitHub.¹

1 INTRODUCTION

Mathematical reasoning is a fundamental domain for assessing the true capabilities of Large Language Models (LLMs) to reason (Ahn et al., 2024). While existing benchmarks like GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021) provide valuable insights, they primarily focus on school-level mathematics. This leaves a significant gap in understanding how LLMs perform on more advanced, university-level problems. Moreover, these benchmarks are becoming saturated, as GPT-4, using advanced prompting techniques, has achieved over 92% success rate on GSM8K and 80% on MATH (Achiam et al., 2023).

Recent works, such as CHAMP (Mao et al., 2024) and MathOdyssey (Fang et al., 2024), aim to introduce more challenging problems but are limited in size (<400 samples) and lack comprehensive topic coverage. The most challenging problems stem from school-level competitions or olympiads, missing the crucial middle ground of university-level coursework that reflects academic demands.

Furthermore, there is a growing interest in assessing multi-modal LLMs' abilities to perform mathematical reasoning involving visual elements (Ahn et al., 2024). Large datasets like MathVista (Lu et al., 2023), We-Math (Qiao et al., 2024), or MathVerse (Zhang et al., 2024) provide an extensive set of (mostly) visual tasks but may lack university-level problems and often rely on multiple-choice validation, leading to easier problems and faster saturation of benchmarks.

*Corresponding author: kchernyshev@toloka.ai

¹<https://github.com/toloka/u-math>

Example: Differential Calculus.

U-MATH Problem:

The function $s(t) = 2 \cdot t^3 - 3 \cdot t^2 - 12 \cdot t + 8$ represents the position of a particle traveling along a horizontal line.

1. Find the velocity and acceleration functions.
2. Determine the time intervals when the object is slowing down or speeding up.

Reference Solution (shortened):

The velocity is $v(t) = s'(t) = 6 \cdot t^2 - 6 \cdot t - 12$, zeros of the $v(t)$ are $t = -1, 2$.

The acceleration is $a(t) = v'(t) = 12 \cdot t - 6$, zero of the $a(t)$ is $t = \frac{1}{2}$.

It speeds up when $v(t)$ and $a(t)$ have the same sign, and slows down when opposite.

Interval	$v(t)$	$a(t)$	Behavior
$(-\infty, -1)$	> 0	< 0	Slowing down
$(-1, \frac{1}{2})$	< 0	< 0	Speeding up
$(\frac{1}{2}, 2)$	< 0	> 0	Slowing down
$(2, \infty)$	> 0	> 0	Speeding up

Accounting for non-negative time, speed up on $(0, 1/2)$ and $(2, \infty)$, slow down on $(1/2, 2)$.

Figure 1: U-MATH covers university-level topics and require multiple steps to solve. A random sample is provided; reference solution is shortened. In this example, common error is overlooking time non-negativity.

In turn, evaluating complex free-form answers remains a significant challenge for the field (Hendrycks et al., 2021). Current methods often rely on LLM judges to assess problems, which introduces potential biases and inconsistencies (Zheng et al., 2023). Errors introduced by automatic evaluators are often overlooked in popular benchmarks. This oversight makes it impossible to account for judge biases, which detracts from the reliability of the evaluation results.

Recent studies also indicate that evaluation of mathematical solutions is a demanding task (Zeng et al., 2023; Xia et al., 2024) and that an LLM’s ability to judge mathematical solutions is correlated with its problem-solving performance (Stephan et al., 2024), further signifying the importance of evaluations designed to assess the evaluators themselves — also called meta-evaluations.

Popular datasets for the task of mathematical meta-evaluation are PRM800K (Lightman et al., 2023), MR-GSM8K (Zeng et al., 2023) and MR-MATH (Xia et al., 2024). However, these are all based on the GSM8K and MATH datasets, still leaving a gap in meta-evaluations for university-level problems.

Aiming to bridge these gaps and provide a comprehensive evaluation of LLMs’ mathematical capabilities, we introduce **U-MATH** (*University Math*) and a supplementary meta-evaluation dataset, which we refer to as **μ -MATH** (*Meta U-MATH*). Our main contributions are:

1. **U-MATH Benchmark** (Section 3): We open-source a set of 1,100 of university-level problems collected from actual coursework with final answers and solutions. About 20% of problems require image *understanding* to be solved. The text-only part of the benchmark is balanced across 6 key subjects: Precalculus, Algebra, Differential Calculus, Integral Calculus, Multivariable Calculus, and Sequences&Series.
2. **μ -MATH Meta-Evaluation Benchmark** (Section 3.3): Additionally, we introduce a set of 1084 meta-evaluation tasks sourced from U-MATH problems and designed to rigorously assess the quality of LLM judges. We manually select approximately 25% of the U-MATH problem statements and golden answers, supplying each with four solutions produced by different top-performing LLMs, and label them based on whether the generated solutions are correct or not. The benchmark is designed to be challenging for LLM judges yet representative of the typical university-level math grading tasks.
3. **Comparison of Models** (Section 4): We conduct a comparative analysis of various open-source and proprietary LLMs on U-MATH. Our analysis highlights the high performance of specialized models in text-only problems and the superiority of proprietary models in visual tasks with the best U-MATH accuracy of 49%. Additionally, we examine several popular

LLMs on μ -MATH to assess their ability to judge free-form mathematical problems. Our results show the best model achieving the macro F1-score of 80%.

We release the U-MATH and μ -MATH benchmarks under a permissive license to facilitate further research and ensure reproducibility.

2 BACKGROUND

Enhancing and evaluating the mathematical reasoning capabilities of LLMs is essential in AI research (Ahn et al., 2024). Studies show that finetuning with mathematical and code-related data enhances models’ general skills (Prakash et al., 2024). Mathematical tasks require logical thinking and multi-step problem-solving, thus improving overall reasoning abilities in LLMs (Chen et al., 2024).

This leads to the problem of evaluating LLM’s math abilities. Despite the significant progress, many existing benchmarks are limited in scope, focusing primarily on school-level mathematics or limited in size and topic coverage. Table 1 summarizes popular text-only and visual mathematical benchmarks.

Dataset	Levels	% Uni. Level	#Test	% Visual	% Free Form Answer
MMLU _{Math} (Hendrycks et al., 2020)	E H C	0	1.3k	0	0
GSM8k (Cobbe et al., 2021)	E	0	1k	0	0
MATH (Hendrycks et al., 2021)	H O	0	5k	0	100
MiniF2F (Zheng et al., 2021)	E H O	0	244	0	100
OCWCourses (Lewkowycz et al., 2022)	U	100	272	0	100
ProofNet (Azerbayev et al., 2023)	C U	≈50	371	0	100
CHAMP (Mao et al., 2024)	H	0	270	0	100
MathOdyssey (Fang et al., 2024)	H U O	26	387	0	100
MMMU _{Math} (Yue et al., 2023)	C	0	505	100	0
MathVista (Lu et al., 2023)	E H C	0	5k	100	46
MATH-V (Wang et al., 2024a)	E H O	0	3k	100	50
We-Math (Qiao et al., 2024)	E H U	≈20	1.7k	100	0
MathVerse (Zhang et al., 2024)	H	0	4.7k	83.3	45
U-MATH (this work)	U	100	1.1k	20	100

Table 1: Existing Auto-evaluation Math benchmarks with corresponding test samples *published*, visual samples percent, and percent of multiple-choice questions. Level denotes E Elementary to Middle School,

Textual Mathematical Benchmarks. Early efforts to assess LLMs’ mathematical abilities have emerged in datasets like MathQA (Amini et al., 2019) and the mathematics subset of MMLU (Hendrycks et al., 2020). These early benchmarks emphasized the importance of operation-based reasoning in solving mathematical word problems, typically in a multiple-choice format. Nowadays, even smaller models (e.g., 7B parameters) have achieved high scores on these tasks (Li et al., 2024b), suggesting that these benchmarks are becoming saturated. In response, more comprehensive datasets have emerged, such as GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021), or MGSM (Shi et al., 2022) (multilingual version of 250 GSM8K samples). These popular benchmarks are crucial for evaluating LLMs’ mathematical reasoning skills. However, they primarily focus on school-level problems, which may not fully assess the depth of mathematical reasoning.

Recent efforts attempt to address more advanced mathematical concepts. MathOdyssey (Fang et al., 2024) with competition problems, OCWCourses (Lewkowycz et al., 2022) from actual MIT courses, and ProofNet (Azerbayev et al., 2023) focusing on proofs aim to evaluate undergraduate-level or olympiad-level knowledge. However, these datasets are constrained by their small sizes (e.g., 387, 272, and 371 samples), limiting their statistical robustness and topic coverage. For example, MathOdyssey is limited to 101 samples in university-level topics (Calculus, Algebra, and Diff. Equations and Statistics). Other specialized datasets like MiniF2F (Zheng et al., 2021) provide valuable parallel corpora in formal languages, while CHAMP (Mao et al., 2024) offers helpful context and hints, but both are similarly limited in scale with 244 and 270 samples. Additionally, both heavily rely on already published resources: CHAMP sources material from a book, while MiniF2F re-uses international olympiads and MATH dataset problems. An attempt to provide a more robust evaluation,

GHOSTS (Frieder et al., 2024) dataset, provides 728 problems (both from other datasets and new ones) but does not provide reference solutions and answers, focusing instead on human evaluation, making cheap automatic evaluation impossible.

The current datasets are either too small, leading to higher measurement errors, or focus mainly on elementary and high school math, leaving a gap in evaluating LLMs’ proficiency in advanced university-level math topics.

Visual Mathematical Benchmarks. As multimodal LLMs gain prominence, there is a growing need for visual mathematical benchmarks (Zhang et al., 2024; Qiao et al., 2024). Early efforts in this domain focus primarily on geometric problems, as seen in datasets like GeoQA (Chen et al., 2022b), UniGeo (Chen et al., 2022a), and Geometry3K (Lu et al., 2021). These datasets have a narrow focus that does not encompass the breadth of mathematical visual reasoning required at advanced levels.

More recent benchmarks attempt to broaden the scope of visual mathematical evaluation. One of the first comprehensive attempts is the mathematical subset of MMMU (Yue et al., 2023), which offers 505 college-level multiple-choice questions, all with images. However, its multiple-choice format limits the complexity of problems that can be posed. MathVista (Lu et al., 2023) collects 28 existing datasets and introduces 3 new datasets with a total of 5k samples (1k testmini samples). However, as shown by Qiao et al. (2024), it faces challenges with data quality due to its compilation from older datasets.

The latest benchmarks, such as MATH-V (Vision) (Wang et al., 2024a) and We-Math (Qiao et al., 2024), extend this approach to collect 3k and 1.7k visual samples, respectively. However, both datasets rely on multiple-choice questions in the test set, leading to faster saturation. MathVerse (Zhang et al., 2024) further extends this approach, relying on visual elements and providing some simple text problems with 1.2k brand-new samples. Among these, only the We-Math dataset includes university-level mathematical problems.

Our U-MATH dataset improves on existing benchmarks with 225 of 1,100 university-level problems that require visual elements (graph, table, diagram) to be solved. This balanced ratio ensures models are challenged to handle both traditional and visual problem-solving without over-relying on visuals, mirroring real-world scenarios.

Large Language Models for Mathematics. The application of LLMs to mathematical problem-solving shows promising results, particularly with models like GPT-3.5 and GPT-4 demonstrating strong reasoning abilities for complex tasks such as those in the MATH dataset (Achiam et al., 2023). While open-source models initially lagged in performance on advanced mathematical tasks, the Llama-3.1 (Dubey et al., 2024) is approaching parity with proprietary models. The most popular benchmarks, MATH and GSM8K, are nearing saturation, with Llama 3.1 405B achieving scores of 73.8% and 96.8%, respectively. Similarly, a Qwen2.5-Math-72B model (Yang et al., 2024b; Team, 2024) reach 85.9% on MATH while Qwen2-Math-72B (Yang et al., 2024a) reaches 96.7% on GSM8k.

To enhance LLMs’ mathematical capabilities, researchers develop various prompt-based methods (Liu et al., 2021). These include techniques for encouraging chain-of-thought generation (Wei et al., 2022), selecting final results from multiple sampled outputs (Wang et al., 2022), and using external tools such as calculators, WolframAlpha or Python interpreters (Gao et al., 2023) to reduce arithmetic errors. Additionally, instruction tuning during pre-training has been identified as a key factor in improving performance (Wang et al., 2017). While these approaches show promise, their effectiveness on university-level problems still needs to be explored due to the lack of suitable large-scale benchmarks.

Mathematical solution verification. Evaluating mathematical solutions is uniquely challenging due to the open-ended nature of answers and the inherent ambiguity in mathematical expressions. Consequently, many benchmarks opt for multiple-choice formats due to their grading simplicity. However, this approach often simplifies tasks, providing hints that models can exploit (Li et al., 2024c; Pezeshkpour and Hruschka, 2023).

While free-form evaluation using LLM judges is widespread (Zheng et al., 2023), it is known to introduce potential errors (Zheng et al., 2023), since evaluating mathematical solutions is a complex task in its own right (Zeng et al., 2023; Xia et al., 2024). These evaluation errors are largely overlooked and unaccounted for, limiting the reliability of inferences drawn from such evaluations.

Hence, it is important to be able to estimate the performance of automatic evaluators and to choose the most adequate among them. Recent studies show that evaluation performance is correlated with but does not equal problem-solving performance (Stephan et al., 2024). This underscores the importance of benchmarks designed specifically to assess the evaluators — also called meta-evaluations.

There are existing benchmarks that are well-suited for meta-evaluations. PRM800K (Lightman et al., 2023) contains 800K annotated steps from 75K solutions to 12K MATH dataset problems, designed to confuse reward models. FELM (Zhao et al., 2024) provides GPT-3.5 annotations for solutions to 208 GSM8K and 194 MATH problems. MR-GSM8K (Zeng et al., 2023) and MR-MATH (Xia et al., 2024) introduce meta-evaluation datasets focused on the GSM8K and MATH datasets, respectively. However, these are either based on elementary to high-school level problems or feature specifically competition-style math, leaving a gap in meta-evaluations on complex and practical university tasks.

To address this, we introduce μ -MATH— a meta-evaluation dataset based on a subset of U-MATH problems. It provides LLM-generated solutions with verified labels, enabling precise and fine-grained assessment of LLMs’ evaluation abilities.

3 U-MATH

We present **U-MATH** (stands for University Math) — a benchmark designed to challenge LLMs with problems requiring deep understanding and advanced reasoning. The problems span 6 core topics and range in difficulty and number of questions. A subset of 20% of problems includes images to test the models’ ability to interpret and reason with graphical information. Reference solutions and answers accompany all problems.

Accuracy is the primary performance metric for **U-MATH**, its text-only problems (**U-MATH_T**) and problems that include a visual component (**U-MATH_V**). The main performance measure for μ -MATH is **macro-F1**.

We use an LLM as a judge (Zheng et al., 2023) to measure the accuracy of the free-form answers against the golden solutions. A problem is considered solved only if all required questions are answered and all requested items (e.g., all saddle points) are correctly identified.

3.1 DATASET COLLECTION

To create a benchmark that authentically reflects university-level mathematics, we collaborate with Gradarius, a platform providing learning content and software for top US universities specialized in mathematics. The problems are sourced from ongoing courses across various institutions currently run on the Gradarius platform. Problems and solutions are crafted by subject matter experts and represent real-world academic standards. These samples are unpublished and have not been exposed to any external sources. Thus, the dataset could not be leaked to current LLMs.

We employ a multi-stage filtering process to select challenging problems from tens of thousands of available samples. First, we filter out problems with short solutions (< 100 characters) and problems in multiple-choice format. As LLMs are not designed to perform arithmetic calculations and are prone to errors (Hendrycks et al., 2021; Lewkowycz et al., 2022), we focus on testing mathematical reasoning rather than calculations. We filter out problems marked as allowing calculator usage. As for the visual problems selection, we chose to keep problems with a single image for convenience.

Next, we employ several small LLMs (LLaMA-3.1-8B (Dubey et al., 2024), Qwen2-7B (Yang et al., 2024a), Mistral-7B (Jiang et al., 2023), Mathstral-7B, NuminaMath-7B (Beeching et al., 2024)) to solve the problems. We select 150 most challenging problems for each subject based on the average problem solution rate. For this step, we use the same pipeline as described in Section 4. This way, we ensure that none of the individual models influence problem selection largely and that there is no overfitting to a specific LLM. As the last step, we hold extra validation high risk problems (with low solve rate) using our in-house math experts and Gradarius content team.

After collection, we enlist a team of experts from the Stevens Institute of Technology, who actively teach various Calculus courses. The experts verify that each problem is suitable either for assessing the subject knowledge expected of college or university students or for testing prerequisite knowledge.

The team thoroughly reviewed and affirmed that the selected problems meet these criteria. Overall, only 4.3% of the problems are categorized as school-level rather than university-level, highlighting the robustness of the selection process.

3.2 DATASET STATISTICS

The U-MATH benchmark comprises **1,100** carefully curated and validated mathematical problems. These problems are distributed across **6 core subjects** with about 20% of the tasks incorporating visual elements, such as graphs, tables, and geometric figures, mirroring the multi-modal nature of real-world mathematical problems: Precalculus (Review), Algebra, Differential Calculus (+Differential Equations), Integral Calculus, Multivariable Calculus, and Sequences & Series.

Math Subject	#Textual	#Visual	Avg. Questions	Avg. Answers
Algebra	150	30	1.93	1.28
Differential Calculus	150	70	2.37	1.15
Integral Calculus	150	58	1.09	1.01
Multivariable Calculus	150	28	1.74	1.09
Precalculus	150	10	1.51	1.23
Sequences and Series	150	4	1.36	1.00
All	900	200	1.66	1.12

Table 2: Average number of questions per problem and answers per question in U-MATH.

Table 2 summarizes the distribution of problems across different subjects. The average is **1.7** questions per problem (e.g., local minima, maxima, and increasing intervals are asked), and the average of **1.1** answers per question (for example, the number of saddle points in the correct answer).

3.3 META-EVALUATION FRAMEWORK (μ -MATH)

The evaluation of mathematical problems is not straightforward. Even simple expressions such as $x \cdot 0.5$ may have valid forms like $\frac{x}{2}$, $x \div 2$, $x/2$, or unsimplified variants like $9x/18$. In practice, evaluating free-form solutions requires testing expression equivalence in much less trivial cases, especially with more advanced problems (refer to Section A.3 in Appendix for an example).

To systematically study the ability of LLMs to evaluate free-form mathematical solutions on advanced, university-level problems, we introduce the μ -MATH (Meta U-MATH) benchmark. It consists of a curated subset of U-MATH samples, supplied with LLM-generated solutions — both correct and not. The solutions are labeled using a combination of manual inspection and automated verification via Gradarius-API, which allows to test formal equivalence of mathematical expressions.

We selected 271 U-MATH problems (around **25%**) based on their assessment difficulty to create a challenging meta-evaluation set. This subset does not aim to reflect the overall U-MATH distribution but rather to provide a robust test for LLM judges. We focused on text-only problems, excluding those needing images, due to the limited size of the labeled U-MATH subset. Four solutions have been generated for each of the selected problems — using Qwen2.5-72B, Llama3.1-8B, GPT-4o and Gemini-1.5-Pro models — **1084 samples** in total.

A tested model is provided with a problem statement, a reference answer, and a solution to evaluate. We treat this as a binary classification task, using the macro-averaged **F1-score as the primary metric** to minimize the effect of class imbalance. Additionally, we report Positive Predictive Value (PPV or Precision) and True Positive Rate (TPR or Recall) for the positive class as well as Negative Predictive Value (NPV) and True Negative Rate (TNR) for the negative class, offering a finer-grained performance evaluation. We also report all of the scores computed both across the entire set of samples and only across those with solutions produced by a specific model, separately for each of the author models.

4 EXPERIMENTS AND RESULTS

4.1 EXPERIMENTAL SETUP

We select some top-performing recent LLMs to evaluate.

Model	Source	Size(s)	Visual	Open-weights
Mathstral-v0.1	(Mistral.ai, 2024)	7B	×	✓
NuminaMath-CoT	(Beeching et al., 2024)	7B	×	✓
LLaMA-3.1	(Dubey et al., 2024)	8B, 70B	×	✓
LLaMA-3.1-Nemotron	(Wang et al., 2024b)	70B	×	✓
Qwen2-Math	(Yang et al., 2024a)	7B, 72B	×	✓
Qwen2.5-Math	(Yang et al., 2024b)	7B, 72B	×	✓
Qwen2.5	(Team, 2024)	7B, 72B	×	✓
Athene-V2-Chat	(Nexusflow, 2024)	72B	×	✓
Pixtral-12B-2409	(Mistral AI, 2024)	12B	✓	✓
LLAVA One Vision _{Qwen2-7B}	(Li et al., 2024a)	8B	✓	✓
Qwen2-VL	(Yang et al., 2024a)	7B, 72B	✓	✓
LLaMA-3.2	(Meta AI, 2024)	11B, 90B	✓	✓
Claude-3.5-Sonnet	(Anthropic, 2024)	unknown	✓	×
GPT-4o-mini-2024-07-18	(OpenAI, 2024)	unknown	✓	×
GPT-4o-2024-08-06	(OpenAI, 2024)	unknown	✓	×
Gemini-1.5-Flash-002	(Team et al., 2024)	unknown	✓	×
Gemini-1.5-Pro-002	(Team et al., 2024)	unknown	✓	×

Table 3: LLMs name, version and sizes we evaluate.

All LLMs are tested using the same prompts and settings for fair comparison. The LLMs are restricted to a single generation of 4096 tokens with the temperature set to 0. We employ chain-of-thought (CoT) prompting (Wei et al., 2022) to encourage models to ‘think’ before providing an answer. Images are included directly in the prompts for multimodal LLMs. To text-only LLMs the problem description is provided as-is without visual elements.

We report accuracy based on GPT-4o-2024-08-06 as-a-judge for our final results, despite it not being the best performing judge — due to the model still residing among the top-ranked judges, being the more conservative one in terms of false positive rate, as well as widely available, leading to easier reproducibility. Details on the judgment setup and comparisons of various judges are all discussed in Section 4.3.

4.2 U-MATH RESULTS

Figure 2 compare popular text-only and multimodal models in U-MATH as well as U-MATH_{Text} and U-MATH_{Visual}. Table 4 summarizes the performance of all evaluated LLMs on the U-MATH benchmark. Reference to Appendix E for model performance vs model size comparison.

Among text-only models, the math-specific model Qwen2.5-Math-72B achieves the highest overall accuracy at 50.2%, showcasing strong mathematical reasoning capabilities. In the multi-modal model group, **Gemini-1.5-pro-002 leads** with an overall accuracy of **60.1%**, highlighting the advantages of integrating visual processing. In contrast, best open-weights model Qwen2-VL-72B lacks mathematical abilities in visual and textual tasks with 31.2% on a U-MATH benchmark. Building on these results, several key trends emerge:

- **Model Size vs. Specialization:** Larger models expectedly outperform smaller ones. However, the small specialized model Qwen2.5-Math-7B surpasses or performs on par with 10 times larger models like Qwen2.5-72B or LLaMA-3.1-70B and almost reaching leading Gemini-1.5-Pro level. On the other hand, Pixtral-12B performs consistently worse than minor Qwen2-VL-7B, indicating a lack of university-level data in training.
- **Textual vs. Visual Problem-Solving:** Across multimodal models, text-only problems’ accuracy vastly exceeds visual problems, highlighting areas for further improvement. The text-only models




Figure 2: Performance of the selected top-performing models on U-MATH, U-MATH_{Text} and U-MATH_{Visual}. Color denotes different model families, ‘visual’ label highlight visual encoder of the model. Higher is better for all charts.

Model	U-MATH	U-MATH		Algebra		Diff. C.		Integral C.		Multivar C.		Precalculus		Seq. & Series	
		T	V	T	V	T	V	T	V	T	V*	T	V*	T	V*
Text-only models															
Mathstral 7B	18.0	20.7	6.0	51.3	6.7	4.0	10.0	1.3	1.7	8.0	3.6	48.7	10.0	10.7	0.0
NuminaMath 7B	19.2	22.8	3.0	62.7	0.0	4.0	7.1	1.3	0.0	6.0	3.6	51.3	0.0	11.3	0.0
Llama-3.1 8B	22.3	26.1	5.0	59.3	3.3	6.7	5.7	9.3	3.4	11.3	3.6	54.7	10.0	15.3	25.0
Qwen2.5 7B	33.8	40.0	6.0	86.0	10.0	12.7	1.4	10.0	12.1	26.7	3.6	75.3	0.0	29.3	0.0
Qwen2.5-Math 7B	38.4	45.2	7.5	87.3	6.7	18.7	5.7	8.0	10.3	36.0	10.7	80.7	0.0	40.7	0.0
NuminaMath 72B	25.0	29.7	4.0	74.7	3.3	6.7	4.3	4.0	3.4	11.3	3.6	62.7	10.0	18.7	0.0
Llama-3.1 70B	28.5	33.7	5.0	82.0	3.3	10.7	5.7	4.0	5.2	14.0	3.6	64.0	0.0	27.3	25.0
Llama-3.1 Nemotron 70B	31.4	37.4	4.0	84.0	0.0	14.7	2.9	4.0	3.4	25.3	7.1	64.0	20.0	32.7	0.0
Qwen2.5 72B	41.0	48.6	7.0	88.7	6.7	22.7	4.3	12.0	6.9	40.0	17.9	83.3	0.0	44.7	0.0
Athene-V2 72B Chat	46.2	54.6	8.5	88.7	3.3	34.0	4.3	16.0	6.9	50.7	21.4	88.7	10.0	49.3	50.0
Qwen2.5-Math 72B	50.2	59.0	10.5	92.7	6.7	35.3	7.1	20.7	17.2	58.0	7.1	90.0	0.0	57.3	50.0
Multimodal models															
Pixtral 12B	15.5	15.6	15.5	44.7	23.3	1.3	34.3	0.7	0.0	3.3	0.0	32.0	0.0	11.3	0.0
Llama-3.2 11B Vision	17.0	18.6	10.0	54.0	10.0	1.3	20.0	1.3	1.7	4.7	3.6	43.3	10.0	6.7	0.0
LLaVA-OV Qwen2-7B	17.7	20.7	4.5	60.7	6.7	4.0	5.7	1.3	1.7	5.3	3.6	43.3	10.0	9.3	0.0
Qwen2-VL 7B	20.4	21.4	15.5	62.7	10.0	4.7	32.9	0.7	5.2	6.7	7.1	45.3	0.0	8.7	0.0
Qwen2-VL 72B	31.2	32.2	26.5	80.7	26.7	9.3	40.0	2.0	13.8	14.7	28.6	65.3	10.0	21.3	0.0
Llama-3.2 90B Vision	32.6	36.3	16.0	85.3	26.7	10.7	25.7	2.7	1.7	22.7	7.1	65.3	20.0	31.3	25.0
Claude Sonnet 3.5	35.1	36.1	30.5	76.0	33.3	12.0	41.4	7.3	17.2	21.3	28.6	65.3	30.0	34.7	25.0
GPT-4o-mini	37.2	40.3	23.0	88.0	16.7	16.7	31.4	4.0	10.3	24.0	35.7	77.3	20.0	32.0	25.0
GPT-4o	43.5	46.4	30.0	91.3	30.0	18.7	32.9	10.0	20.7	41.3	42.9	79.3	30.0	38.0	25.0
Gemini 1.5 Flash	51.3	53.8	40.0	91.3	50.0	36.0	45.7	14.0	24.1	44.0	50.0	80.7	30.0	56.7	50.0
Gemini 1.5 Pro	60.1	63.4	45.0	91.3	60.0	50.7	47.1	27.3	24.1	60.7	57.1	87.3	70.0	63.3	50.0

Table 4: Comparison of models’ accuracy on our U-MATH benchmark and its subjects. Scores for various mathematical categories, including text and visual analysis, are displayed. For each subject 2 numbers are provided - text-only (T) and visual (V) problems. Asterisk denotes a small number of samples (< 30). Free-form solutions judged by gpt-4o-2024-08-06. Images are not included in the prompt for text-only models, only the problem statement. **Bold** indicates the best result in each group.

can solve a small percentage of visual problems, primarily due to guessing or judgment errors discussed in Section 4.3.

- **Proprietary vs. Open-weights model:** Proprietary models like Gemini still offer top or competitive performance but lack transparency and flexibility. At the moment, the gap is evident in visual comprehension, with 18.5% difference on U-MATH_{Visual} between top-1 and best open-weight model. However, open-weight models like Qwen-Math is a big step toward top performance.
- **Continuous Finetuning:** Additional tuning significantly enhances performance, with LLaMA-3.1 70B \Rightarrow LLaMA-3.1 Nemotron 70B and Qwen2.5-72B \Rightarrow Athene-V2 72B achieving 2.9% and 5.2% higher U-MATH accuracy, respectively. This reinforces the idea that models are not fully optimized for their size and require high-quality data for further improvements.

Subject-Specific Results Model performance varies across mathematical subjects, excelling in text-based tasks for Precalculus and Algebra, consistent with benchmark saturation (Ahn et al., 2024), but faltering on visual-symbolic tasks. In Sequences and Series, success on formula-based

problems reflects logical structuring, though limited visual data restricts evaluation. Differential and Multivariable Calculus results are moderate, with difficulties in abstract, multi-dimensional concepts, especially visual interpretations. Integral Calculus presents the greatest challenge, as interpreting curves, areas, and extensive expressions confounds models, underscoring the need for improved multimodal training.

4.3 META-EVALUATION (μ -MATH) RESULTS

For meta-evaluation we use the same setup as described in Section 4.1. Additionally, we experiment with two distinct prompting schemes — a standard Automatic Chain-of-Thought (AutoCoT) prompt involving a simple task description together with an instruction to think step-by-step, and a manual Chain-of-Thought prompt (which we refer to as simply CoT) with explicit instructions on which steps to follow when approaching the task. We find the latter prompting scheme to perform best, so we use manual CoT for the main results. The judge’s output is also further processed by an extractor model (Qwen2.5 72B is fixed for consistency), prompted to produce a single label — either ‘Yes’, ‘No’ or ‘Inconclusive’. We include ‘Inconclusive’ for cases when judge refuses to evaluate or generation fails; such judgments are treated as incorrect. Reference Appendix C.2 for full contents of all the prompts.

Model	U-MATH _{Test}	μ -MATH					μ -MATH _{Qwen}	μ -MATH _{Llama}	μ -MATH _{GPT}	μ -MATH _{Gemini}
		F1 _{CoT} / F1 _{AutoCoT}	TPR	TNR	PPV	NPV	F1 _{CoT} / F1 _{AutoCoT}	F1 _{CoT} / F1 _{AutoCoT}	F1 _{CoT} / F1 _{AutoCoT}	F1 _{CoT} / F1 _{AutoCoT}
Llama-3.1 8B	26.1	52.0 / 53.1	48.7	55.9	56.0	48.5	48.7 / 49.6	49.2 / 51.2	51.2 / 57.6	55.5 / 50.5
Qwen2.5 7B	40.0	69.3 / 67.0	78.7	59.8	69.3	70.8	62.4 / 60.5	72.3 / 72.4	68.3 / 66.4	69.1 / 65.0
Qwen2.5-Math 7B	45.2	61.9 / 61.2	76.6	47.9	62.9	63.9	59.7 / 56.7	63.8 / 64.0	57.2 / 58.5	63.8 / 61.2
GPT-4o-mini	40.3	72.3 / 69.2	59.0	88.1	85.1	65.1	69.3 / 61.7	76.2 / 78.5	70.4 / 69.8	69.6 / 64.3
Gemini 1.5 Flash	53.8	74.8 / 65.3	63.3	88.3	86.2	67.6	71.2 / 61.9	80.6 / 70.8	70.1 / 65.3	73.9 / 59.7
Llama-3.1-70B	33.7	61.0 / 68.2	62.5	59.6	64.1	57.9	56.0 / 63.8	57.0 / 70.2	69.4 / 69.8	58.8 / 64.4
Qwen2.5 72B	48.6	75.6 / 75.1	77.1	74.2	77.5	73.7	70.5 / 68.9	79.3 / 80.1	73.7 / 73.4	74.2 / 73.8
Qwen2.5-Math 72B	59.0	74.0 / 75.5	80.9	66.8	73.8	75.2	69.3 / 68.8	77.3 / 79.8	68.2 / 69.2	76.8 / 80.4
Claude 3.5 Sonnet	36.1	74.8 / 68.1	62.5	89.5	87.3	67.4	70.8 / 64.1	77.9 / 71.8	72.2 / 68.1	73.8 / 63.4
GPT-4o	46.4	77.4 / 74.2	70.1	85.9	85.1	71.3	74.2 / 68.2	81.8 / 78.9	77.5 / 75.8	72.6 / 70.5
Gemini 1.5 Pro	63.4	80.7 / 69.1	77.5	84.5	85.2	76.4	77.7 / 64.6	83.6 / 74.8	78.2 / 68.6	79.5 / 64.9

Table 5: Comparison of model’s ability to judge on μ -MATH benchmark using CoT prompting; Macro F1-score (F1), True Positive Rate (TPR), True Negative Rate (TNR), Positive Predictive Value (PPV) and Negative Predictive Value (NPV) are presented, with F1 as the primary one. The second number within each F1 column written in gray represents the F1-score under AutoCoT prompting. μ -MATH columns represent integral scores over the entire benchmark, while μ -MATH_{<model>} columns denote subsets with solutions generated by specific author models. U-MATH_{Text} accuracy is added for comparison of each model’s performance as a math solver vs as a math judge. **Bold** indicates the best result in each column. Full expanded tables are presented in Appendix K.

We find that **using manual CoT instructions instead of the standard AutoCoT improves or maintains judgment performance**, save for Llama models, as shown in Table 5. Llama’s performance drop is largely due to increased inconclusive judgment rates (see Appendix G). At the same time, Gemini models benefit the most from this transition, gaining over 10% in F1-score and becoming the top-ranked models, surpassing Qwen and GPT models that outperform Gemini in the AutoCoT setting. This shows that **prompting effects are substantial yet inhomogeneous across models**. Please refer to Appendix F for a visual comparison.

In terms of the resulting performance, we see that **correctly identifying a positive label is harder on average** compared to negative labels, with the best TPR being almost 10% lower than the best TNR, and that **the best attainable F1 score is only 80.7%**. This constitutes a considerable deficiency in the context of judgment, because judges’ error rates directly limit the precision of capability evaluations, potentially even biasing them in case the errors are systematic in nature as opposed to pure noise.

Our results, for instance, reveal a **consistent bias towards some models** — better performance on Llama solutions and worse performance on Qwen solutions — most pronounced with smaller-sized judges and AutoCoT prompting. This bias is generally reduced for both small and large judges when transitioning to CoT prompting, which is also illustrated with Figure 3. At the same time, no noticeable ‘self-judgment’ effects are found.

It is also evident that **being a better solver does not necessarily lead to being a better judge**. In fact, our results suggest a trade-off existing between these skill; refer to Appendix H for visualizations and a more detailed discussion.




Figure 3: Relative differences between specific judgment performance — i.e. over samples with solutions generated by a specific author model — and integral judgment performance across all the samples. The judgment performance is measured by the μ -MATH macro F1-scores. Each pane corresponds to a different author model considered when measuring specific performance. The x-axis specifies which judge corresponds to a particular bar pair, with bar pairs comparing the above-described relative diffs in case of AutoCoT and CoT prompting schemes.

Besides that, we observe substantive differences in judges’ behavior: **proprietary models tend to be more conservative** — having relatively high TNR compared to their TPR — while **Qwen family of models exhibits the opposite pattern**. The behavior differences are further studied and illustrated in Appendix I.

Overall, judges show varying behaviors, they have imperfect performance that is also distinct from problem solving performance, and different prompting schemes induce nontrivial changes in judges’ behaviors, biases, and even their performance rankings.

All of these findings underscore the importance of performing meta-evaluations, since such things are impossible to quantify and comparisons impossible to make in the absence of datasets designed to benchmark the judges.

5 CONCLUSION

We introduce **U-MATH**, a novel multimodal benchmark for evaluating the university-level mathematical reasoning of LLMs. U-MATH includes 1,100 unpublished free-form problems from real teaching materials, covering 6 core mathematical subjects, with 20% involving image-based reasoning. Additionally, we provide **μ -MATH**, a meta-evaluation dataset, to assesses LLMs’ ability to evaluate free-form mathematical solutions.

Our experiments highlight significant challenges for LLMs in advanced reasoning and visual problem-solving. The highest accuracy achieved was 63.4% on text-based tasks and 45.0% on visual problems (Gemini-1.5-pro-002). Solution assessment remains difficult, with Gemini-1.5-pro-002 achieving top μ -MATH F1-score of 80%, showing room for improvement and underscoring the limitations of widely used models like GPT-4o in evaluation tasks.

Limitations. While U-MATH offers diverse university-level problems, it does not cover the full range of advanced topics and may introduce biases by favoring certain problem types. Also, selection process may introduce biases, potentially favoring certain problem types or difficulty levels (e.g., more accessible topics like Precalculus and Algebra). The inclusion of 20% visual problems, yet reflect real distribution, limits the evaluation of visual reasoning. Furthermore, reliance on LLMs for valuation introduces potential, as models struggle with complex reasoning and instructions, evidenced by our findings with the μ -MATH. The μ -MATH dataset encompass of 25% of U-MATH problems narrows the evaluation scope, but provide 4 diverse model families as solution generators.

Future Work. Future research can focus on enhancing LLM performance by integrating existing tool-augmented models and exploring their effectiveness on U-MATH and μ -MATH tasks. For instance, incorporating external tools, such as formal solvers, could improve complex textual and multimodal reasoning capabilities. Additionally, our findings indicate that widely used models like GPT-4o are not a silver bullet for solution evaluation; thus, developing specialized (finetuned) models or techniques for more accurate and unbiased assessment is a promising direction. Expanding μ -

MATH with formal verification methods could further enhance the evaluation processes. Additionally, conducting deeper prompt sensitivity analyses would provide valuable insights for the field.

By open-sourcing U-MATH, μ -MATH, and the evaluation code, we aim to facilitate further research in advancing the mathematical reasoning capabilities of LLMs and encourage the development of models better equipped to tackle complex, real-world mathematical problems.

ACKNOWLEDGEMENT

We thank all contributors from Stevens Institute and Toloka.ai experts who assisted in sourcing and verifying problems, inspect the solutions and provided valuable feedback throughout the development of U-MATH.

We would like to give special thanks to the dedicated team of Gradarius experts who played a crucial role in validating problems and ensuring their quality:

Jan Cannizzo	Paul Schwartz
Andrey Nikolaev	Arina Voorhaar
Chloe Weiers	Funda Gul
Tatiana Zenkevich	Igor Teplukhovskiy
Anastasiia Feklina	Ruslan Akmetdinov
Ekaterina Eremina	Sofia Tekhazheva

ETHICS STATEMENT

We collected all data in U-MATH and μ -MATH with appropriate permissions, ensuring no personal or proprietary information is included. The datasets consist solely of mathematical problems and solutions, without any sensitive content. We open-sourced the datasets and code under suitable licenses to support transparency and research advancement. There are no known conflicts of interest associated with this work.

REPRODUCIBILITY STATEMENT

All datasets and code will be available on GitHub. Detailed descriptions of dataset collection and processing are in Section 3. The experimental setup, including model configurations and prompts, is outlined in Section 4, with full prompts provided in Appendices C.1 and C.2. These resources enable replication of our experiments.

REFERENCES

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Janice Ahn, Rishu Verma, Renze Lou, Di Liu, Rui Zhang, and Wenpeng Yin. 2024. Large language models for mathematical reasoning: Progresses and challenges.
- Aida Amini, Saadia Gabriel, Shanchuan Lin, Rik Koncel-Kedziorski, Yejin Choi, and Hannaneh Hajishirzi. 2019. MathQA: Towards interpretable math word problem solving with operation-based formalisms. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2357–2367, Minneapolis, Minnesota. Association for Computational Linguistics.

-
- Anthropic. 2024. Introducing claude 3.5 sonnet. <https://www.anthropic.com/news/claude-3-5-sonnet>. Accessed: 2024-11-20.
- Zhangir Azerbayev, Bartosz Piotrowski, Hailey Schoelkopf, Edward W Ayers, Dragomir Radev, and Jeremy Avigad. 2023. Proofnet: Autoformalizing and formally proving undergraduate-level mathematics. *arXiv preprint arXiv:2302.12433*.
- Edward Beeching, Shengyi Costa Huang, Albert Jiang, Jia Li, Benjamin Lipkin, Zihan Qina, Kashif Rasul, Ziju Shen, Roman Soletskyi, and Lewis Tunstall. 2024. NuminaMath 7b cot. <https://huggingface.co/AI-MO/NuminaMath-7B-CoT>.
- Jiaqi Chen, Tong Li, Jinghui Qin, Pan Lu, Liang Lin, Chongyu Chen, and Xiaodan Liang. 2022a. UniGeo: Unifying geometry logical reasoning via reformulating mathematical expression. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3313–3323, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Jiaqi Chen, Jianheng Tang, Jinghui Qin, Xiaodan Liang, Lingbo Liu, Eric P. Xing, and Liang Lin. 2022b. Geoqa: A geometric question answering benchmark towards multimodal numerical reasoning.
- Nuo Chen, Ning Wu, Jianhui Chang, and Jia Li. 2024. Controlmath: Controllable data generation promotes math generalist models.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Meng Fang, Xiangpeng Wan, Fei Lu, Fei Xing, and Kai Zou. 2024. Mathodyssey: Benchmarking mathematical problem-solving skills in large language models using odyssey math data. *arXiv preprint arXiv:2406.18321*.
- Simon Frieder, Luca Pinchetti, Ryan-Rhys Griffiths, Tommaso Salvatori, Thomas Lukasiewicz, Philipp Petersen, and Julius Berner. 2024. Mathematical capabilities of chatgpt. *Advances in neural information processing systems*, 36.
- Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. 2023. Pal: Program-aided language models. In *International Conference on Machine Learning*, pages 10764–10799. PMLR.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2020. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*.
- Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2023. Mistral 7b.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. 2022. Solving quantitative reasoning problems with language models.
- Bo Li, Yuanhan Zhang, Dong Guo, Renrui Zhang, Feng Li, Hao Zhang, Kaichen Zhang, Yanwei Li, Ziwei Liu, and Chunyuan Li. 2024a. Llava-onevision: Easy visual task transfer.

-
- Chen Li, Weiqi Wang, Jingcheng Hu, Yixuan Wei, Nanning Zheng, Han Hu, Zheng Zhang, and Houwen Peng. 2024b. Common 7b language models already possess strong math capabilities.
- Wangyue Li, Liangzhi Li, Tong Xiang, Xiao Liu, Wei Deng, and Noa Garcia. 2024c. Can multiple-choice questions really be useful in detecting the abilities of llms?
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2023. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*.
- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2021. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing.
- Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. 2023. Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts. *arXiv preprint arXiv:2310.02255*.
- Pan Lu, Ran Gong, Shibiao Jiang, Liang Qiu, Siyuan Huang, Xiaodan Liang, and Song-Chun Zhu. 2021. Inter-gps: Interpretable geometry problem solving with formal language and symbolic reasoning.
- Yujun Mao, Yoon Kim, and Yilun Zhou. 2024. Champ: A competition-level dataset for fine-grained analyses of llms’ mathematical reasoning capabilities. *arXiv preprint arXiv:2401.06961*.
- Meta AI. 2024. Llama 3.2: Revolutionizing edge ai and vision with open, customizable models. <https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/>. Accessed: 2024-11-15.
- Mistral AI. 2024. Announcing pixtral-12b. <https://mistral.ai/news/pixtral-12b/>. Accessed: 2024-10-01.
- Mistral.ai. 2024. Mathstral. <https://mistral.ai/news/mathstral/>. Accessed: 2024-10-01.
- Nexusflow. 2024. Introducing athene-v2: Advancing beyond the limits of scaling with targeted post-training. <https://nexusflow.ai/blogs/athene-v2>. Accessed: 2024-11-15.
- OpenAI. 2024. Hello gpt-4o. <https://openai.com/index/hello-gpt-4o/>. Accessed: 2024-10-01.
- Pouya Pezeshkpour and Estevam Hruschka. 2023. Large language models sensitivity to the order of options in multiple-choice questions.
- Nikhil Prakash, Tamar Rott Shaham, Tal Haklay, Yonatan Belinkov, and David Bau. 2024. Fine-tuning enhances existing mechanisms: A case study on entity tracking.
- Runqi Qiao, Qiuna Tan, Guanting Dong, Minhui Wu, Chong Sun, Xiaoshuai Song, Zhuoma GongQue, Shanglin Lei, Zhe Wei, Miaoquan Zhang, et al. 2024. We-math: Does your large multimodal model achieve human-like mathematical reasoning? *arXiv preprint arXiv:2407.01284*.
- Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Soroush Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, et al. 2022. Language models are multilingual chain-of-thought reasoners. *arXiv preprint arXiv:2210.03057*.
- Andreas Stephan, Dawei Zhu, Matthias Aßenmacher, Xiaoyu Shen, and Benjamin Roth. 2024. From calculation to adjudication: Examining llm judges on mathematical reasoning tasks.
- Gemini Team, Petko Georgiev, Ving Ian Lei, Ryan Burnell, Libin Bai, Anmol Gulati, Garrett Tanzer, Damien Vincent, Zhufeng Pan, Shibo Wang, Soroosh Mariooryad, Yifan Ding, Xinyang Geng, Fred Alcober, Roy Frostig, Mark Omernick, Lexi Walker, Cosmin Paduraru, Christina Sorokin, Andrea Tacchetti, Colin Gaffney, Samira Daruki, Olcan Sercinoglu, Zach Gleicher, Juliette Love, Paul Voigtlaender, Rohan Jain, Gabriela Surita, Kareem Mohamed, Rory Blevins, Junwhan Ahn, Tao Zhu, Kornrhapop Kawintiranon, Orhan Firat, Yiming Gu, Yujing Zhang, Matthew Rahtz, Manaal Faruqui, Natalie Clay, Justin Gilmer, JD Co-Reyes, Ivo Penchev, Rui Zhu, Nobuyuki Morioka,

Kevin Hui, Krishna Haridasan, Victor Campos, Mahdis Mahdiah, Mandy Guo, Samer Hassan, Kevin Kilgour, Arpi Vezar, Heng-Tze Cheng, Raoul de Liedekerke, Siddharth Goyal, Paul Barham, DJ Strouse, Seb Noury, Jonas Adler, Mukund Sundararajan, Sharad Vikram, Dmitry Lepikhin, Michela Paganini, Xavier Garcia, Fan Yang, Dasha Valter, Maja Trebacz, Kiran Vodrahalli, Chulayuth Asawaroengchai, Roman Ring, Norbert Kalb, Livio Baldini Soares, Siddhartha Brahma, David Steiner, Tianhe Yu, Fabian Mentzer, Antoine He, Lucas Gonzalez, Bibo Xu, Raphael Lopez Kaufman, Laurent El Shafey, Junhyuk Oh, Tom Hennigan, George van den Driessche, Seth Odoom, Mario Lucic, Becca Roelofs, Sid Lall, Amit Marathe, Betty Chan, Santiago Ontanon, Luheng He, Denis Teplyashin, Jonathan Lai, Phil Crone, Bogdan Damoc, Lewis Ho, Sebastian Riedel, Karel Lenc, Chih-Kuan Yeh, Aakanksha Chowdhery, Yang Xu, Mehran Kazemi, Ehsan Amid, Anastasia Petrushkina, Kevin Swersky, Ali Khodaei, Gowoon Chen, Chris Larkin, Mario Pinto, Geng Yan, Adria Puigdomenech Badia, Piyush Patil, Steven Hansen, Dave Orr, Sebastien M. R. Arnold, Jordan Grimstad, Andrew Dai, Sholto Douglas, Rishika Sinha, Vikas Yadav, Xi Chen, Elena Gribovskaya, Jacob Austin, Jeffrey Zhao, Kaushal Patel, Paul Komarek, Sophia Austin, Sebastian Borgeaud, Linda Friso, Abhimanyu Goyal, Ben Caine, Kris Cao, Da-Woon Chung, Matthew Lamm, Gabe Barth-Maron, Thais Kagohara, Kate Olszewska, Mia Chen, Kaushik Shivakumar, Rishabh Agarwal, Harshal Godhia, Ravi Rajwar, Javier Snider, Xerxes Dotiwalla, Yuan Liu, Aditya Barua, Victor Ungureanu, Yuan Zhang, Bat-Orgil Batsaikhan, Mateo Wirth, James Qin, Ivo Danihelka, Tulsee Doshi, Martin Chadwick, Jilin Chen, Sanil Jain, Quoc Le, Arjun Kar, Madhu Gurumurthy, Cheng Li, Ruoxin Sang, Fangyu Liu, Lampros Lamprou, Rich Munoz, Nathan Lintz, Harsh Mehta, Heidi Howard, Malcolm Reynolds, Lora Aroyo, Quan Wang, Lorenzo Blanco, Albin Cassirer, Jordan Griffith, Dipanjan Das, Stephan Lee, Jakub Sygnowski, Zach Fisher, James Besley, Richard Powell, Zafarali Ahmed, Dominik Paulus, David Reitter, Zalan Borsos, Rishabh Joshi, Aedan Pope, Steven Hand, Vittorio Selo, Vihan Jain, Nikhil Sethi, Megha Goel, Takaki Makino, Rhys May, Zhen Yang, Johan Schalkwyk, Christina Butterfield, Anja Hauth, Alex Goldin, Will Hawkins, Evan Senter, Sergey Brin, Oliver Woodman, Marvin Ritter, Eric Noland, Minh Giang, Vijay Bolina, Lisa Lee, Tim Blyth, Ian Mackinnon, Machel Reid, Obaid Sarvana, David Silver, Alexander Chen, Lily Wang, Loren Maggiore, Oscar Chang, Nithya Attaluri, Gregory Thornton, Chung-Cheng Chiu, Oskar Bunyan, Nir Levine, Timothy Chung, Evgenii Eltyshev, Xiance Si, Timothy Lillicrap, Demetra Brady, Vaibhav Aggarwal, Boxi Wu, Yuanzhong Xu, Ross McIlroy, Kartikeya Badola, Paramjit Sandhu, Erica Moreira, Wojciech Stokowiec, Ross Hemsley, Dong Li, Alex Tudor, Pranav Shyam, Elahe Rahimtoroghi, Salem Haykal, Pablo Sprechmann, Xiang Zhou, Diana Mincu, Yujia Li, Ravi Addanki, Kalpesh Krishna, Xiao Wu, Alexandre Frechette, Matan Eyal, Allan Dafoe, Dave Lacey, Jay Whang, Thi Avrahami, Ye Zhang, Emanuel Taropa, Hanzhao Lin, Daniel Toyama, Eliza Rutherford, Motoki Sano, HyunJeong Choe, Alex Tomala, Chalence Safranek-Shrader, Nora Kassner, Mantas Pajarskas, Matt Harvey, Sean Sechrist, Meire Fortunato, Christina Lyu, Gamaleldin Elsayed, Chenkai Kuang, James Lottes, Eric Chu, Chao Jia, Chih-Wei Chen, Peter Humphreys, Kate Baumli, Connie Tao, Rajkumar Samuel, Cicero Nogueira dos Santos, Anders Andreassen, Nemanja Rakićević, Dominik Grewe, Aviral Kumar, Stephanie Winkler, Jonathan Caton, Andrew Brock, Sid Dalmia, Hannah Sheahan, Iain Barr, Yingjie Miao, Paul Natsev, Jacob Devlin, Feryal Behbahani, Flavien Prost, Yanhua Sun, Artiom Myaskovsky, Thanumalayan Sankaranarayanan Pillai, Dan Hurt, Angeliki Lazaridou, Xi Xiong, Ce Zheng, Fabio Pardo, Xiaowei Li, Dan Horgan, Joe Stanton, Moran Ambar, Fei Xia, Alejandro Lince, Mingqiu Wang, Basil Mustafa, Albert Webson, Hyo Lee, Rohan Anil, Martin Wicke, Timothy Dozat, Abhishek Sinha, Enrique Piqueras, Elahe Dabir, Shyam Upadhyay, Anudhyan Boral, Lisa Anne Hendricks, Corey Fry, Josip Djolonga, Yi Su, Jake Walker, Jane Labanowski, Ronny Huang, Vedant Misra, Jeremy Chen, RJ Skerry-Ryan, Avi Singh, Shruti Rijhwani, Dian Yu, Alex Castro-Ros, Beer Changpinyo, Romina Datta, Sumit Bagri, Arnar Mar Hrafnkelsson, Marcello Maggioni, Daniel Zheng, Yury Sulsky, Shaobo Hou, Tom Le Paine, Antoine Yang, Jason Riesa, Dominika Rogozinska, Dror Marcus, Dalia El Badawy, Qiao Zhang, Luyu Wang, Helen Miller, Jeremy Greer, Lars Lowe Sjoes, Azade Nova, Heiga Zen, Rahma Chaabouni, Mihaela Rosca, Jiepu Jiang, Charlie Chen, Ruibo Liu, Tara Sainath, Maxim Krikun, Alex Polozov, Jean-Baptiste Lespiau, Josh Newlan, Zeynep Cankara, Soo Kwak, Yunhan Xu, Phil Chen, Andy Coenen, Clemens Meyer, Katerina Tsihlias, Ada Ma, Juraj Gottweis, Jinwei Xing, Chenjie Gu, Jin Miao, Christian Frank, Zeynep Cankara, Sanjay Ganapathy, Ishita Dasgupta, Steph Hughes-Fitt, Heng Chen, David Reid, Keran Rong, Hongmin Fan, Joost van Amersfoort, Vincent Zhuang, Aaron Cohen, Shixiang Shane Gu, Anhad Mohananey, Anastasija Ilic, Taylor Tobin, John Wieting, Anna Bortsova, Phoebe Thacker, Emma Wang, Emily Caveness, Justin Chiu, Eren Sezener, Alex Kaskasoli, Steven Baker, Katie Millican, Mohamed Elhawaty, Kostas Aisopos, Carl Lebsack, Nathan Byrd,

Hanjun Dai, Wenhao Jia, Matthew Wiethoff, Elnaz Davoodi, Albert Weston, Lakshman Yagati, Arun Ahuja, Isabel Gao, Golan Pundak, Susan Zhang, Michael Azzam, Khe Chai Sim, Sergi Caelles, James Keeling, Abhanshu Sharma, Andy Swing, YaGuang Li, Chenxi Liu, Carrie Grimes Bostock, Yamini Bansal, Zachary Nado, Ankesh Anand, Josh Lipschultz, Abhijit Karmarkar, Lev Proleev, Abe Ittycheriah, Soheil Hassas Yeganeh, George Polovets, Aleksandra Faust, Jiao Sun, Alban Rrustemi, Pen Li, Rakesh Shivanna, Jeremiah Liu, Chris Welty, Federico Lebron, Anirudh Baddepudi, Sebastian Krause, Emilio Parisotto, Radu Soricut, Zheng Xu, Dawn Bloxwich, Melvin Johnson, Behnam Neyshabur, Justin Mao-Jones, Renshen Wang, Vinay Ramasesh, Zaheer Abbas, Arthur Guez, Constant Segal, Duc Dung Nguyen, James Svensson, Le Hou, Sarah York, Kieran Milan, Sophie Bridgers, Wiktor Gworek, Marco Tagliasacchi, James Lee-Thorp, Michael Chang, Alexey Guseynov, Ale Jakse Hartman, Michael Kwong, Ruizhe Zhao, Sheleem Kashem, Elizabeth Cole, Antoine Miech, Richard Tanburn, Mary Phuong, Filip Pavetic, Sebastien Cevey, Ramona Comanescu, Richard Ives, Sherry Yang, Cosmo Du, Bo Li, Zizhao Zhang, Mariko Iinuma, Clara Huiyi Hu, Aurko Roy, Shaan Bijwadia, Zhenkai Zhu, Danilo Martins, Rachel Saputro, Anita Gergely, Steven Zheng, Dawei Jia, Ioannis Antonoglou, Adam Sadovsky, Shane Gu, Yingying Bi, Alek Andreev, Sina Samangooei, Mina Khan, Tomas Kocisky, Angelos Filos, Chintu Kumar, Colton Bishop, Adams Yu, Sarah Hodgkinson, Sid Mittal, Premal Shah, Alexandre Moufarek, Yong Cheng, Adam Bloniarz, Jaehoon Lee, Pedram Pejman, Paul Michel, Stephen Spencer, Vladimir Feinberg, Xuehan Xiong, Nikolay Savinov, Charlotte Smith, Siamak Shakeri, Dustin Tran, Mary Chesus, Bernd Bohnet, George Tucker, Tamara von Glehn, Carrie Muir, Yiran Mao, Hideto Kazawa, Ambrose Slone, Kedar Soparkar, Disha Shrivastava, James Cobon-Kerr, Michael Sharman, Jay Pavagadhi, Carlos Araya, Karolis Misiunas, Nimesh Ghelani, Michael Laskin, David Barker, Qiujia Li, Anton Briukhov, Neil Houlsby, Mia Glaese, Balaji Lakshminarayanan, Nathan Schucher, Yunhao Tang, Eli Collins, Hyeontaek Lim, Fangxiaoyu Feng, Adria Recasens, Guangda Lai, Alberto Magni, Nicola De Cao, Aditya Siddhant, Zoe Ashwood, Jordi Orbay, Mostafa Dehghani, Jenny Brennan, Yifan He, Kelvin Xu, Yang Gao, Carl Saroufim, James Molloy, Xinyi Wu, Seb Arnold, Solomon Chang, Julian Schrittwieser, Elena Buchatskaya, Soroush Radpour, Martin Polacek, Skye Giordano, Ankur Bapna, Simon Tokumine, Vincent Hellendoorn, Thibault Sottiaux, Sarah Cogan, Aliaksei Severyn, Mohammad Saleh, Shantanu Thakoor, Laurent Shefey, Siyuan Qiao, Meenu Gaba, Shuo yin Chang, Craig Swanson, Biao Zhang, Benjamin Lee, Paul Kishan Rubenstein, Gan Song, Tom Kwiatkowski, Anna Koop, Ajay Kannan, David Kao, Parker Schuh, Axel Stjerngren, Golnaz Ghiasi, Gena Gibson, Luke Vilnis, Ye Yuan, Felipe Tiengo Ferreira, Aishwarya Kamath, Ted Klimenko, Ken Franko, Kefan Xiao, Indro Bhattacharya, Miteyan Patel, Rui Wang, Alex Morris, Robin Strudel, Vivek Sharma, Peter Choy, Sayed Hadi Hashemi, Jessica Landon, Mara Finkelstein, Priya Jhakra, Justin Frye, Megan Barnes, Matthew Mauger, Dennis Daun, Khuslen Baatarsukh, Matthew Tung, Wael Farhan, Henryk Michalewski, Fabio Viola, Felix de Chaumont Quitry, Charline Le Lan, Tom Hudson, Qingze Wang, Felix Fischer, Ivy Zheng, Elspeth White, Anca Dragan, Jean baptiste Alayrac, Eric Ni, Alexander Pritzel, Adam Iwanicki, Michael Isard, Anna Bulanova, Lukas Zilka, Ethan Dyer, Devendra Sachan, Srivatsan Srinivasan, Hannah Muckenhirn, Honglong Cai, Amol Mandhane, Mukarram Tariq, Jack W. Rae, Gary Wang, Kareem Ayoub, Nicholas FitzGerald, Yao Zhao, Woohyun Han, Chris Alberti, Dan Garrette, Kashyap Krishnakumar, Mai Gimenez, Anselm Levskaya, Daniel Sohn, Josip Matak, Inaki Iturrate, Michael B. Chang, Jackie Xiang, Yuan Cao, Nishant Ranka, Geoff Brown, Adrian Hutter, Vahab Mirrokni, Nanxin Chen, Kaisheng Yao, Zoltan Egyed, Francois Galilee, Tyler Liechty, Praveen Kallakuri, Evan Palmer, Sanjay Ghemawat, Jasmine Liu, David Tao, Chloe Thornton, Tim Green, Mimi Jasarevic, Sharon Lin, Victor Cotruta, Yi-Xuan Tan, Noah Fiedel, Hongkun Yu, Ed Chi, Alexander Neitz, Jens Heitkaemper, Anu Sinha, Denny Zhou, Yi Sun, Charbel Kaed, Brice Hulse, Swaroop Mishra, Maria Georgaki, Sneha Kudugunta, Clement Farabet, Izhak Shafran, Daniel Vlasic, Anton Tsitsulin, Rajagopal Ananthanarayanan, Alen Carin, Guolong Su, Pei Sun, Shashank V, Gabriel Carvajal, Josef Broder, Iulia Comsa, Alena Repina, William Wong, Warren Weilun Chen, Peter Hawkins, Egor Filonov, Lucia Loher, Christoph Hirsenschall, Weiye Wang, Jingchen Ye, Andrea Burns, Hardie Cate, Diana Gage Wright, Federico Piccinini, Lei Zhang, Chu-Cheng Lin, Ionel Gog, Yana Kulizhskaya, Ashwin Sreevatsa, Shuang Song, Luis C. Cobo, Anand Iyer, Chetan Tekur, Guillermo Garrido, Zhu Yun Xiao, Rupert Kemp, Huaixiu Steven Zheng, Hui Li, Ananth Agarwal, Christel Ngani, Kati Goshvadi, Rebeca Santamaria-Fernandez, Wojciech Fica, Xinyun Chen, Chris Gorgolewski, Sean Sun, Roopal Garg, Xinyu Ye, S. M. Ali Eslami, Nan Hua, Jon Simon, Pratik Joshi, Yelin Kim, Ian Tenney, Sahitya Potluri, Lam Nguyen Thiet, Quan Yuan, Florian Luisier, Alexandra Chronopoulou, Salvatore Scellato, Praveen Srinivasan, Minmin Chen, Vinod Koverkathu, Valentin Dalibard, Yaming Xu, Brennan Saeta, Keith Anderson, Thibault Sellam,

Nick Fernando, Fantine Huot, Junehyuk Jung, Mani Varadarajan, Michael Quinn, Amit Raul, Maigo Le, Ruslan Habalov, Jon Clark, Komal Jalan, Kalesha Bullard, Achintya Singhal, Thang Luong, Boyu Wang, Sujeevan Rajayogam, Julian Eischenschlos, Johnson Jia, Daniel Finchelstein, Alex Yakubovich, Daniel Balle, Michael Fink, Sameer Agarwal, Jing Li, Dj Dvijotham, Shalini Pal, Kai Kang, Jaclyn Konzelmann, Jennifer Beattie, Olivier Dousse, Diane Wu, Remi Crocker, Chen Elkind, Siddhartha Reddy Jonnalagadda, Jong Lee, Dan Holtmann-Rice, Krystal Kallarackal, Rosanne Liu, Denis Vnukov, Neera Vats, Luca Invernizzi, Mohsen Jafari, Huanjie Zhou, Lilly Taylor, Jennifer Prendki, Marcus Wu, Tom Eccles, Tianqi Liu, Kavya Kopparapu, Francoise Beaufays, Christof Angermueller, Andreea Marzoca, Shourya Sarcar, Hilal Dib, Jeff Stanway, Frank Perbet, Nejc Trdin, Rachel Sterneck, Andrey Khorlin, Dinghua Li, Xihui Wu, Sonam Goenka, David Madras, Sasha Goldshtein, Willi Gierke, Tong Zhou, Yaxin Liu, Yannie Liang, Anais White, Yunjie Li, Shreya Singh, Sanaz Bahargam, Mark Epstein, Sujoy Basu, Li Lao, Adnan Ozturk, Carl Crous, Alex Zhai, Han Lu, Zora Tung, Neeraj Gaur, Alanna Walton, Lucas Dixon, Ming Zhang, Amir Globerson, Grant Uy, Andrew Bolt, Olivia Wiles, Milad Nasr, Ilia Shumailov, Marco Selvi, Francesco Piccinno, Ricardo Aguilar, Sara McCarthy, Misha Khalman, Mrinal Shukla, Vlado Galic, John Carpenter, Kevin Vilella, Haibin Zhang, Harry Richardson, James Martens, Matko Bosnjak, Shreyas Rammohan Belle, Jeff Seibert, Mahmoud Alnahlawi, Brian McWilliams, Sankalp Singh, Annie Louis, Wen Ding, Dan Popovici, Lenin Simicich, Laura Knight, Pulkit Mehta, Nishesh Gupta, Chongyang Shi, Saaber Fatehi, Jovana Mitrovic, Alex Grills, Joseph Pagadora, Dessie Petrova, Danielle Eisenbud, Zhishuai Zhang, Damion Yates, Bhavishya Mittal, Nilesh Tripuraneni, Yannis Assael, Thomas Brovelli, Prateek Jain, Mihajlo Velimirovic, Canfer Akbulut, Jiaqi Mu, Wolfgang Macherey, Ravin Kumar, Jun Xu, Haroon Qureshi, Gheorghe Comanici, Jeremy Wiesner, Zhitao Gong, Anton Ruddock, Matthias Bauer, Nick Felt, Anirudh GP, Anurag Arnab, Dustin Zelle, Jonas Rothfuss, Bill Rosgen, Ashish Shenoy, Bryan Seybold, Xinjian Li, Jayaram Mudigonda, Goker Erdogan, Jiawei Xia, Jiri Simsa, Andrea Michi, Yi Yao, Christopher Yew, Steven Kan, Isaac Caswell, Carey Radebaugh, Andre Elisseeff, Pedro Valenzuela, Kay McKinney, Kim Paterson, Albert Cui, Eri Latorre-Chimoto, Solomon Kim, William Zeng, Ken Durden, Priya Ponnappalli, Tiberiu Sosea, Christopher A. Choquette-Choo, James Manyika, Brona Robenek, Harsha Vashisht, Sebastien Pereira, Hoi Lam, Marko Velic, Denese Owusu-Afriyie, Katherine Lee, Tolga Bolukbasi, Alicia Parrish, Shawn Lu, Jane Park, Balaji Venkatraman, Alice Talbert, Lambert Rosique, Yuchung Cheng, Andrei Sozanschi, Adam Paszke, Praveen Kumar, Jessica Austin, Lu Li, Khalid Salama, Wooyeol Kim, Nandita Dukkupati, Anthony Baryshnikov, Christos Kaplanis, XiangHai Sheng, Yuri Chervonyi, Caglar Unlu, Diego de Las Casas, Harry Askham, Kathryn Tunyasuvunakool, Felix Gimeno, Siim Poder, Chester Kwak, Matt Miecznikowski, Vahab Mirrokni, Alek Dimitriev, Aaron Parisi, Dangyi Liu, Tomy Tsai, Toby Shevlane, Christina Kouridi, Drew Garmon, Adrian Goedeckemeyer, Adam R. Brown, Anitha Vijayakumar, Ali Elqursh, Sadegh Jazayeri, Jin Huang, Sara Mc Carthy, Jay Hoover, Lucy Kim, Sandeep Kumar, Wei Chen, Courtney Biles, Garrett Bingham, Evan Rosen, Lisa Wang, Qijun Tan, David Engel, Francesco Pongetti, Dario de Cesare, Dongseong Hwang, Lily Yu, Jennifer Pullman, Srini Narayanan, Kyle Levin, Siddharth Gopal, Megan Li, Asaf Aharoni, Trieu Trinh, Jessica Lo, Norman Casagrande, Roopali Vij, Loic Matthey, Bramandia Ramadhana, Austin Matthews, CJ Carey, Matthew Johnson, Kremena Goranova, Rohin Shah, Shereen Ashraf, Kingshuk Dasgupta, Rasmus Larsen, Yicheng Wang, Manish Reddy Vuyyuru, Chong Jiang, Joana Ijazi, Kazuki Osawa, Celine Smith, Ramya Sree Boppana, Taylan Bilal, Yuma Koizumi, Ying Xu, Yasemin Altun, Nir Shabat, Ben Bariach, Alex Korchemniy, Kiam Choo, Olaf Ronneberger, Chimezie Iwuanyanwu, Shubin Zhao, David Soergel, Cho-Jui Hsieh, Irene Cai, Shariq Iqbal, Martin Sundermeyer, Zhe Chen, Elie Bursztein, Chaitanya Malaviya, Fadi Biadisy, Prakash Shroff, Inderjit Dhillon, Tejasi Latkar, Chris Dyer, Hannah Forbes, Massimo Nicosia, Vitaly Nikolaev, Somer Greene, Marin Georgiev, Pidong Wang, Nina Martin, Hanie Sedghi, John Zhang, Praseem Banzal, Doug Fritz, Vikram Rao, Xuezhi Wang, Jiageng Zhang, Viorica Patraucean, Dayou Du, Igor Mordatch, Ivan Jurin, Lewis Liu, Ayush Dubey, Abhi Mohan, Janek Nowakowski, Vlad-Doru Ion, Nan Wei, Reiko Tojo, Maria Abi Raad, Drew A. Hudson, Vaishakh Keshava, Shubham Agrawal, Kevin Ramirez, Zhichun Wu, Hoang Nguyen, Ji Liu, Madhavi Sewak, Bryce Petrini, DongHyun Choi, Ivan Philips, Ziyue Wang, Ioana Bica, Ankush Garg, Jarek Wilkiewicz, Priyanka Agrawal, Xiaowei Li, Danhao Guo, Emily Xue, Naseer Shaik, Andrew Leach, Sadh MNM Khan, Julia Wiesinger, Sammy Jerome, Abhishek Chakladar, Alek Wenjiao Wang, Tina Ornduff, Folake Abu, Alireza Ghaffarkhah, Marcus Wainwright, Mario Cortes, Frederick Liu, Joshua Maynez, Andreas Terzis, Pouya Samangouei, Riham Mansour, Tomasz Kępa, François-Xavier Aubet, Anton Algymr, Dan Banica, Agoston Weisz, Andras Orban, Alexandre Senges, Ewa Andrejczuk, Mark Geller, Niccolo Dal Santo,


-
- Valentin Anklin, Majd Al Merey, Martin Baeuml, Trevor Strohmaier, Junwen Bai, Slav Petrov, Yonghui Wu, Demis Hassabis, Koray Kavukcuoglu, Jeffrey Dean, and Oriol Vinyals. 2024. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context.
- Qwen Team. 2024. Qwen2.5: A party of foundation models.
- Ke Wang, Junting Pan, Weikang Shi, Zimu Lu, Mingjie Zhan, and Hongsheng Li. 2024a. Measuring multimodal mathematical reasoning with math-vision dataset. *arXiv preprint arXiv:2402.14804*.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2022. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*.
- Yan Wang, Xiaojiang Liu, and Shuming Shi. 2017. Deep neural solver for math word problems. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 845–854, Copenhagen, Denmark. Association for Computational Linguistics.
- Zhilin Wang, Alexander Bukharin, Olivier Delalleau, Daniel Egert, Gerald Shen, Jiaqi Zeng, Oleksii Kuchaiev, and Yi Dong. 2024b. Helpsteer2-preference: Complementing ratings with preferences.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837.
- Shijie Xia, Xuefeng Li, Yixin Liu, Tongshuang Wu, and Pengfei Liu. 2024. Evaluating mathematical reasoning beyond accuracy. *arXiv preprint arXiv:2404.05692*.
- An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, Guanting Dong, Haoran Wei, Huan Lin, Jialong Tang, Jialin Wang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Ma, Jianxin Yang, Jin Xu, Jingren Zhou, Jinze Bai, Jinzheng He, Junyang Lin, Kai Dang, Keming Lu, Keqin Chen, Kexin Yang, Mei Li, Mingfeng Xue, Na Ni, Pei Zhang, Peng Wang, Ru Peng, Rui Men, Ruize Gao, Runji Lin, Shijie Wang, Shuai Bai, Sinan Tan, Tianhang Zhu, Tianhao Li, Tianyu Liu, Wenbin Ge, Xiaodong Deng, Xiaohuan Zhou, Xingzhang Ren, Xinyu Zhang, Xipin Wei, Xuancheng Ren, Xuejing Liu, Yang Fan, Yang Yao, Yichang Zhang, Yu Wan, Yunfei Chu, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, Zhifang Guo, and Zhihao Fan. 2024a. Qwen2 technical report.
- An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu, Xingzhang Ren, and Zhenru Zhang. 2024b. Qwen2.5-math technical report: Toward mathematical expert model via self-improvement.
- Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruoqi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, et al. 2023. Mmmu: A massive multi-discipline multimodal understanding and reasoning benchmark for expert agi. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9556–9567.
- Zhongshen Zeng, Pengguang Chen, Shu Liu, Haiyun Jiang, and Jiaya Jia. 2023. Mr-gsm8k: A meta-reasoning benchmark for large language model evaluation. *CoRR*, abs/2312.17080.
- Renrui Zhang, Dongzhi Jiang, Yichi Zhang, Haokun Lin, Ziyu Guo, Pengshuo Qiu, Aojun Zhou, Pan Lu, Kai-Wei Chang, Peng Gao, et al. 2024. Mathverse: Does your multi-modal llm truly see the diagrams in visual math problems? *arXiv preprint arXiv:2403.14624*.
- Yiran Zhao, Jinghan Zhang, I Chern, Siyang Gao, Pengfei Liu, Junxian He, et al. 2024. Felm: Benchmarking factuality evaluation of large language models. *Advances in Neural Information Processing Systems*, 36.
- Kunhao Zheng, Jesse Michael Han, and Stanislas Polu. 2021. Minif2f: a cross-system benchmark for formal olympiad-level mathematics. *arXiv preprint arXiv:2109.00110*.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena.

A PROBLEM EXAMPLES

A.1 U-MATH PROBLEMS

Example 1: Algebra.

Write a logarithmic equation corresponding to the graph shown. Use $\log_3(x)$ as a parent function:



The final answer: $-3 \cdot \log_3(x + 4)$

Example 2: Integral Calculus.

Solve the integral:

$$\int \frac{-9 \cdot \sqrt[3]{x}}{9 \cdot \sqrt[3]{x^2} + 3 \cdot \sqrt{x}} dx$$


$$= -\frac{2}{27} \cdot \ln \left(\frac{1}{3} \cdot |1 + 3 \cdot \sqrt[6]{x}| \right)$$

$$= -\frac{1}{3} \cdot \sqrt[6]{x^2} - \frac{3}{2} \cdot \sqrt[6]{x^4} + \frac{2}{3} \cdot \sqrt[6]{x^3}$$

$$+ \frac{2}{9} \cdot \sqrt[6]{x} + C$$

Example 3: Precalculus Review.

Find a formula for $f(x)$, the sinusoidal function whose graph is shown below:



The final answer:
 $f(x) = -4 \cdot \cos \left(\frac{1}{2} \cdot \left(x - \frac{\pi}{4} \right) \right) + 4$


Example 4: Multivariable Calculus.

E is located inside the cylinder $x^2 + y^2 = 1$ and between the circular paraboloids $z = 1 - x^2 - y^2$ and $z = x^2 + y^2$. Find the volume of E .

Volume = $\frac{\pi}{4}$

Example 5: Multivariable Calculus.

The graph of the polar rectangular region D is given. Express the region D in polar coordinates:



1. The interval of r is $[3, 5]$
2. The interval of θ is $\left[\frac{3}{4} \cdot \pi, \frac{5}{4} \cdot \pi \right]$

Example 6: Differential Calculus.

Sketch the curve:

$$y = \frac{x^3}{6 \cdot (x + 3)^2}$$

Provide the following:

1. The domain (in interval notation)
2. Vertical asymptotes
3. Horizontal asymptotes
4. Slant asymptotes
5. Intervals where the function is increasing
6. Intervals where the function is decreasing
7. Intervals where the function is concave up
8. Intervals where the function is concave down
9. Points of inflection

1. The domain: $(-\infty, -3) \cup (-3, \infty)$

Figure 4: Example text-only and visual problems from the U-MATH benchmark, illustrating the topic, problem, and golden answer.

A.2 U-MATH PROBLEM AND SOLUTION

Example: Differential Calculus.

U-MATH Problem:

The function $s(t) = 2 \cdot t^3 - 3 \cdot t^2 - 12 \cdot t + 8$ represents the position of a particle traveling along a horizontal line.

1. Find the velocity and acceleration functions.
2. Determine the time intervals when the object is slowing down or speeding up.

Golden answer:

1. The velocity function $v(t) = 6 \cdot t^2 - 6 \cdot t - 12$ and acceleration function $a(t) = 12 \cdot t - 6$
2. The time intervals when the object speeds up $(0, \frac{1}{2})$, $(2, \infty)$ and slows down $(\frac{1}{2}, 2)$

Reference solution:

The velocity is the derivative of the position function:

$$v(t) = s'(t) = 6t^2 - 6t - 12$$

The acceleration is the derivative of the velocity function:

$$a(t) = v'(t) = 12t - 6$$

To determine when the object is speeding up or slowing down, we compare the signs of $v(t)$ and $a(t)$.

Step 1: Find the Zeros of $v(t)$ and $a(t)$

First, solve for $v(t) = 0$:

$$6t^2 - 6t - 12 = 0 \Rightarrow t^2 - t - 2 = 0 \Rightarrow (t - 2)(t + 1) = 0$$

Thus, $t = 2$ and $t = -1$.

Next, solve for $a(t) = 0$:

$$12t - 6 = 0 \Rightarrow t = \frac{1}{2}$$

Step 2: Analyze the Signs of $v(t)$ and $a(t)$

We analyze the signs of $v(t)$ and $a(t)$ on the intervals determined by $t = -1$, $t = \frac{1}{2}$, and $t = 2$.

Interval	$v(t)$	$a(t)$	Behavior
$(-\infty, -1)$	> 0	< 0	Slowing down
$(-1, \frac{1}{2})$	< 0	< 0	Speeding up
$(\frac{1}{2}, 2)$	< 0	> 0	Slowing down
$(2, \infty)$	> 0	> 0	Speeding up

Step 3: Account for non-negative time

The object is speeding up on $(0, \frac{1}{2})$ and $(2, \infty)$ and slowing down on $(\frac{1}{2}, 2)$.

Figure 5: An example problem from the U-MATH benchmark, illustrating the problem, reference solution and golden answer.

A.3 μ -MATH META-EVALUATION

Example: Integral Calculus.

U-MATH Problem:
Solve the integral:

$$\int \frac{20 \cdot \cos(-10 \cdot x)^3}{21 \cdot \sin(-10 \cdot x)^7} dx$$

Golden answer:

$$C + \frac{1}{21} \cdot \left(\frac{1}{2} \cdot (\cot(10 \cdot x))^4 + \frac{1}{3} \cdot (\cot(10 \cdot x))^6 \right)$$

LLM-generated answer:

$$-\frac{3 \sin(10x)^2 - 2}{126 \sin(10x)^6} + C$$

Golden judge verdict: correct

Comment:
Omitting the arbitrary constants, the reference and the submission could be expressed, respectively, as

$$\frac{\cot^4(10x)}{42} + \frac{\cot^6(10x)}{63} \quad \text{and} \quad \frac{\cot^6(10x)}{63} + \frac{\cot^4(10x)}{42} + \frac{1}{126},$$

which differ by a constant term of 1/126.

Figure 6: An example problem from the μ -MATH meta-evaluation benchmark, illustrating the comparison between the golden (reference) answer and the answer generated by an LLM.

B SUB-TOPICS DISTRIBUTION

The U-MATH dataset cover variety of topics across 6 core subjects. Below is the count of unique topics per subject:

- Differential Calculus: 51 unique topics
- Sequences and Series: 28 unique topics
- Integral Calculus: 35 unique topics
- Precalculus Review: 19 unique topics
- Algebra: 74 unique topics
- Multivariable Calculus: 53 unique topics

Subject	Topic Count	Topic Name
Differential Calculus	29	Curve Sketching
	13	Limits
	12	One-Sided Limits
	12	L'Hospital's Rule
	11	Increasing and Decreasing Functions
	11	Higher Derivatives
	10	Applications of Derivatives (Local Extrema)
Sequences and Series	40	Taylor Series
	30	Fourier Series
	18	Maclaurin Series
	12	Approximating Constants Using Power Series
	6	Radius of Convergence (Center of Convergence)
	5	Differentiate Power Series
	4	Error in Approximation
Integral Calculus	83	The Substitution Rule
	24	Antiderivatives
	10	Volumes of Solids of Revolution About the X-Axis
	9	Trigonometric Substitutions and Inverse Substitutions
	9	Integrate Respect Independent Variable
	7	Applications of Integrals
Precalculus Review	7	Single Variable Surface Area Integrals
	55	Trigonometric Functions
	24	Zeros
	11	Inverses of Functions
	8	Inequalities
	7	Equations with Exponents and Logarithms
	7	Properties of Functions
Algebra	6	Exponential Functions
	18	Equations and Inequalities
	13	Polynomial Equations
	8	Find Composition of Two Functions
	7	Polynomials
	6	Find Slope Line
	6	Applications of Exponential Function
Multivariable Calculus	6	Quadratic Equations
	13	Triple Integrals
	11	Lagrange Multipliers
	9	Double Integrals in Polar Coordinates
	8	Derivatives of Parametric Equations
	8	Integrals of Multivariable Functions
	8	Double Integral Over General Region
	6	Classification of Critical Points

Table 6: Top 7 Topics for Each Subject.

C PROMPTS

C.1 PREDICTION PROMPT

Solution CoT Prompt.

```
{{problem}}\nPlease reason step by step, and put your final answer within \boxed{ }
```

Comment:
Images (if present) are passed with native for provider API schema. For OpenAI-compatible endpoints it is image_url field.^a

^a<https://platform.openai.com/docs/guides/vision>

Figure 7: Prediction for comparing student’s answer and reference answer

C.2 JUDGMENT PROMPTS

Judgment Automatic CoT Prompt.

You'll be provided with a math problem, a correct answer for it and a solution for evaluation.
You have to answer whether the solution is correct or not.

PROBLEM STATEMENT:

{{problem}}

CORRECT ANSWER:

{{golden_answer}}

SOLUTION TO EVALUATE:

{{generated_solution}}

Now please compare the answer obtained in the solution with the provided correct answer to evaluate whether the solution is correct or not.

Think step-by-step, then conclude with your final verdict by putting either "Yes" or "No" on a separate line.

Figure 8: Judgment Automatic CoT Prompt for comparing student's answer and reference answer. This prompt has not been used in U-MATH evaluation.

Judgment CoT Prompt.

You'll be provided with a math problem, a correct answer for it and a solution for evaluation.
You have to answer whether the solution is correct or not.

PROBLEM STATEMENT:

{{problem}}

CORRECT ANSWER:

{{golden_answer}}

SOLUTION TO EVALUATE:

{{generated_solution}}

Now please compare the answer obtained in the solution with the provided correct answer to evaluate whether the solution is correct or not.

Think step-by-step, following these steps, don't skip any:

1. Extract the answer from the provided solution
2. Make any derivations or transformations that may be necessary to compare the provided correct answer with the extracted answer
3. Perform the comparison
4. Conclude with your final verdict — put either "Yes" or "No" on a separate line

Figure 9: Judgment CoT Prompt for comparing student's answer and reference answer. This is the prompt that has been used in U-MATH evaluation.

Judgment Extract Prompt.

You'll be given a result of an evaluation of some mathematical solution by a professional evaluator. You need to extract the final verdict of this evaluation in simple terms: is the solution graded as correct or not.

Output only a single label — "Yes", "No" or "Inconclusive" — according to the provided evaluation ("Yes" if the solution is graded as correct, "No" if the solution is graded as incorrect, "Inconclusive" if the evaluation is incomplete or the final verdict is not settled upon).

Only output "Inconclusive" for incomplete or unsettled evaluations. If the evaluation does contain a single final verdict like "Yes", "Correct", "True", "No", "Incorrect", "False" and so on, even if it is supplied with some additional disclaimers and remarks, output a "Yes" or "No" label accordingly.

Here goes your input:

'''

{{generated_judgment}}

'''

Now please output exactly either "Yes", "No" or "Inconclusive".

Figure 10: Prompt for extracting the final verdict from the judge's outputs.

D SOLUTION PREDICTIONS LENGTH DISTRIBUTION

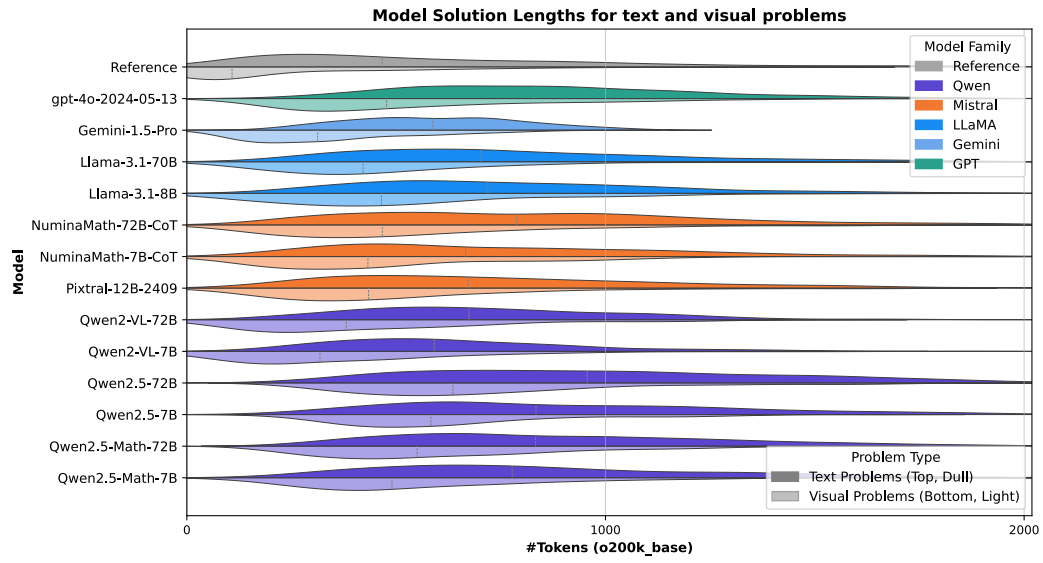


Figure 11: Distribution of token number for generated solutions: Text-only problems (top, dull) and Visual problems (bottom, light). o200k_base tokenizer is used for consistency.

E MODEL ACCURACY VS SIZE

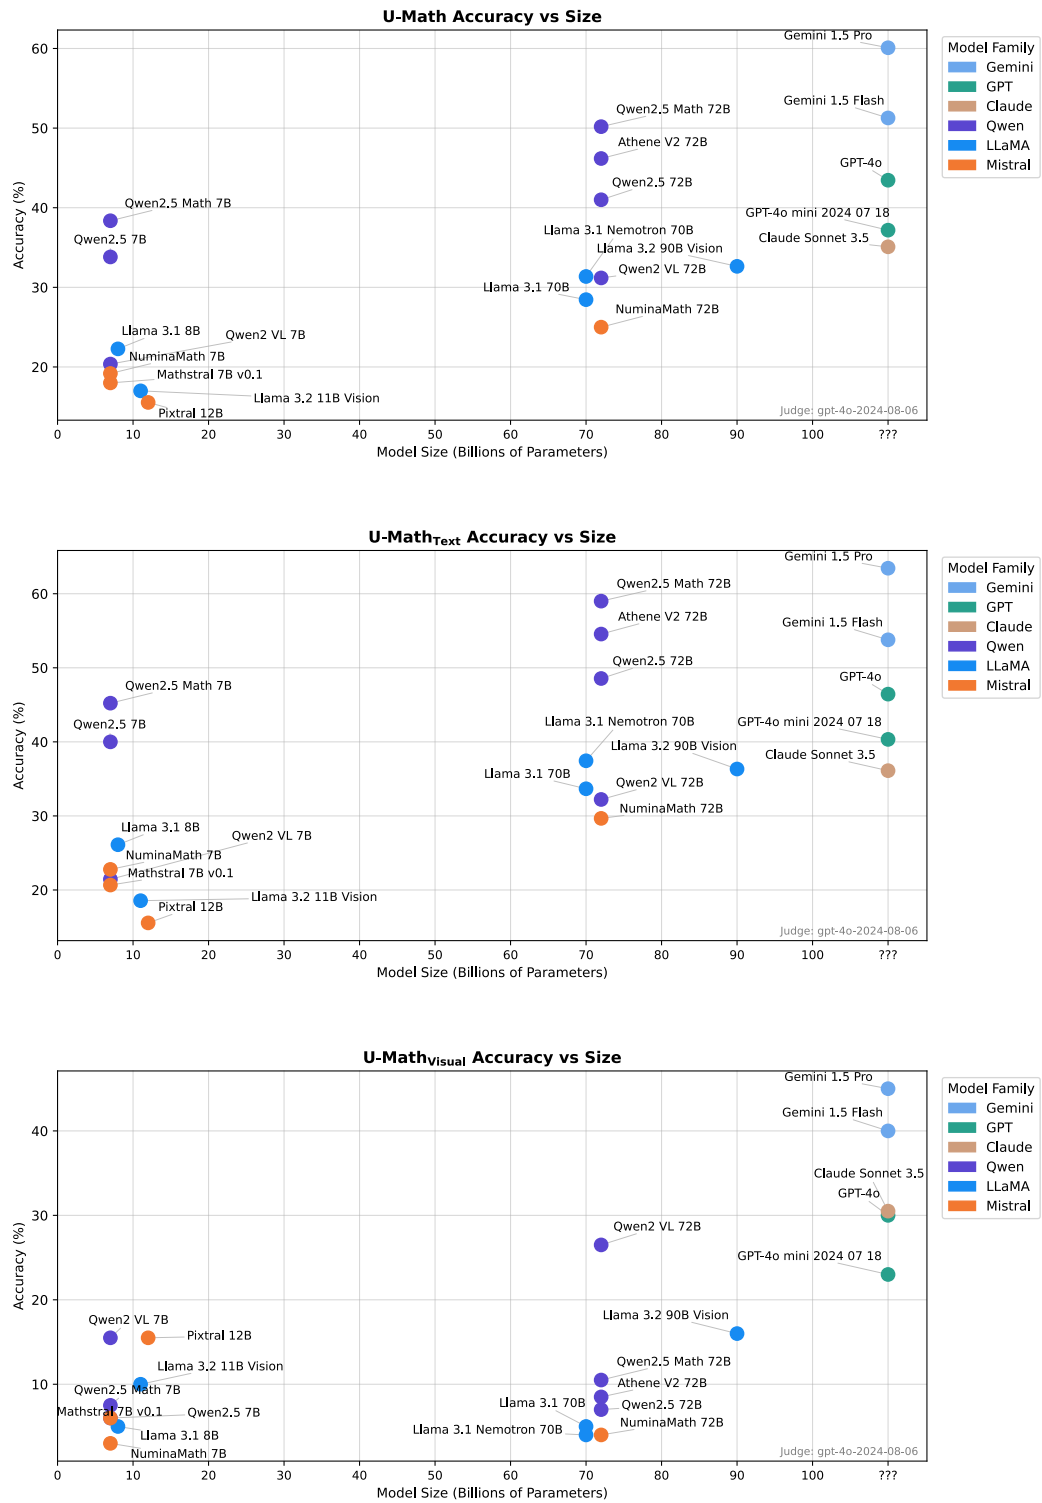


Figure 12: Accuracy of the selected top-performing models on U-MATH, U-MATH_{Text}, and U-MATH_{Visual}. Color denotes different model families. Higher is better for all charts.

F μ -MATH PROMPTING SCHEMES COMPARISON

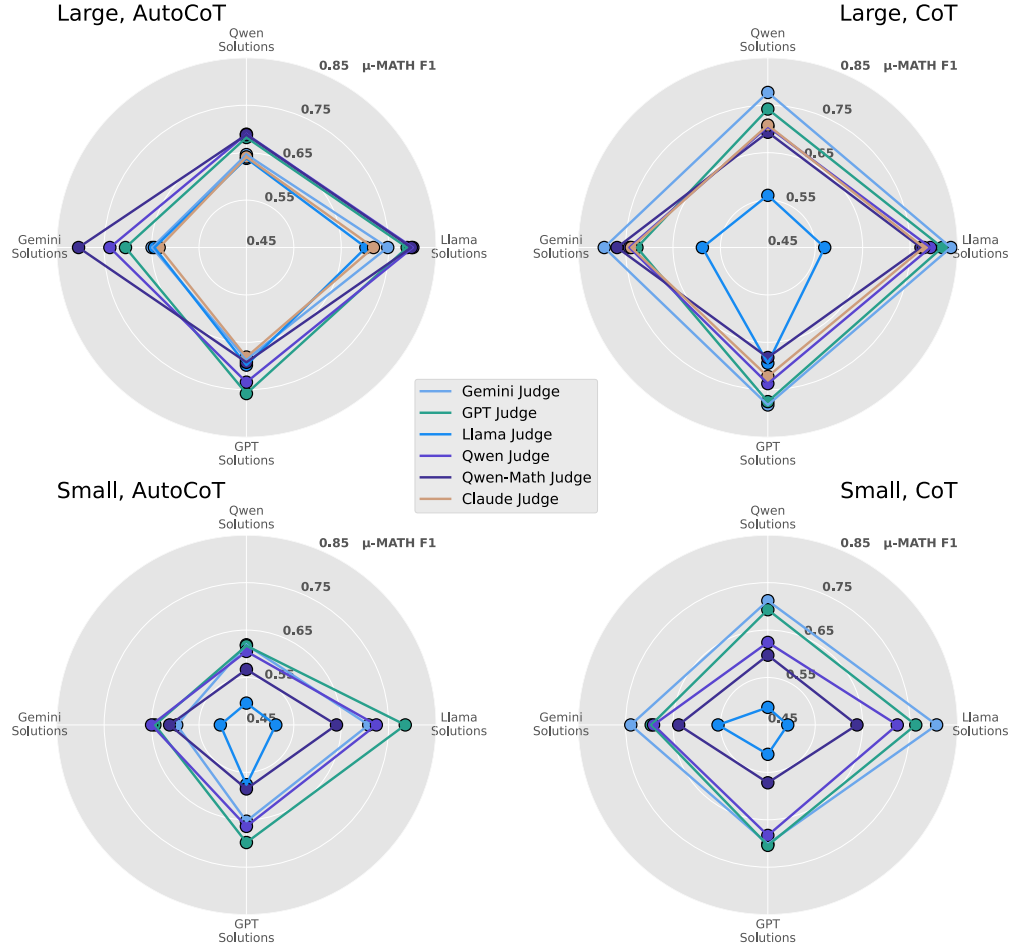


Figure 13: μ -MATH macro F1-score for each of the four solution author models, split by the judge model size and prompting scheme used.

G μ -MATH INCONCLUSIVE JUDGMENT RATES

	Llama-3.1 8B	Llama-3.1 70B	Qwen2.5-Math 7B	Qwen2.5-Math 72B	Qwen2.5 7B	Qwen2.5 72B	GPT-4o-mini	GPT-4o	Gemini 1.5 Flash	Gemini 1.5 Pro	Claude 3.5 Sonnet
AutoCoT	13.4	5.0	2.8	1.2	1.0	1.6	0.0	0.0	0.0	0.0	0.0
CoT	22.9	13.8	2.4	0.7	1.2	2.1	0.1	0.0	0.1	0.0	0.0

Table 7: Percentages of inconclusive judgments produced by each model under different prompting schemes.

H COMPARISON OF PROBLEM SOLVING AND JUDGMENT PERFORMANCE

In this section, we provide a detailed comparison of models’ performances on U-MATH and μ -MATH. The overall distribution of scores visualized in Figure 14 not only shows that improved problem-solving performance does not immediately lead to better judgment performance, as discussed in Section 4.3, but also **suggests a possible trade-off** present between these capabilities.

This possibility is further illustrated when considering specific models. For instance, Qwen2.5-Math demonstrates strong problem solving compared to most of the other models, but does so at the expense of weaker instruction following — eye-gaze inspections reveal this model to be struggling with instruction comprehension and adherence to formatting rules — leading to lower relative judgment performance. In contrast, Claude does not rank low as a judge despite its weak performance on U-MATH. Meanwhile, Gemini, known to excel in both mathematical problem solving and instruction following, comes out as the top-ranked judge.




Figure 14: Comparison of LLMs’ problem solving (U-MATH_{Text}) vs judgment (μ -MATH) performance.

I μ -MATH BEHAVIOR OF JUDGES

In Figure 15 we visualize the difference in ‘performance profiles’ of the judges which we’ve discussed in Section 4.3 — proprietary models behaving more conservatively and Qwen family models exhibiting the opposite tendencies.

This is in line with eye-gazing inspections suggesting that:

- **Qwens tend to ‘follow the solution’** and are also good at going into involved derivation chains necessary to arrive at a true positive verdict in more complex scenarios, albeit **at the cost of increased hallucination risk**.
- **Proprietary models are more ‘anchored on the label’** and less heavy on long hallucination-prone transformation chains, which comes **at the expense of missing more complex true positives**.

Notably, **Claude Sonnet 3.5 and Qwen2.5-Math 72B are the ‘opposite extremes’** of the observed patterns — having respectively the highest overall TNR and highest overall TPR — with an approximately equal F1-score. To illustrate the patterns, Appendix J provides an example comparing the Claude’s and Qwen’s judgments on a single μ -MATH sample. Notice how Claude is restrictive and superficial in its comparison, whereas Qwen ‘loses the structure’ along the way, designating only the first two steps prescribed with the CoT prompt (see prompt contents in Appendix C.2) and omitting points three and four.

These effects are further amplified with reduction in model size: proprietary models mainly lose in TPR when moving from larger models to smaller ones, whereas Qwens, once again on the contrary, lose more in TNR. A possible interpretation is that model size helps to mitigate the natural tendencies potentially induced by the training data, perhaps due to better generalization across general and domains-specific skills, or due to increased reliability.

Also, among the higher-ranking judges, **the large Qwen2.5 and Gemini models turn out to be the two most balanced representatives** of their respective source classes. We speculate that these two models are trained with data explicitly balancing math-specific skills and general capabilities: Gemini being heavily optimized for both math and instruction following, and Qwen likely trained on a blend containing synthetic data produced by various specialist models, including mathematical synthetic data from Qwen2-Math.

The behavioral differences may also be observed with predicted label agreement rates between judges, see Figure 16 for the comparison. Interestingly, **no pair of models has agreement above around 80%** — even for same-family models like Qwen2.5 and Qwen2.5-Math — despite the pairwise μ -MATH performance deltas being small compared to 20% disagreement.

All of this shows that judge comparison is substantive beyond the one-dimensional choice of the better model and suggests judge ensembling to be a potentially fruitful approach to evaluation.

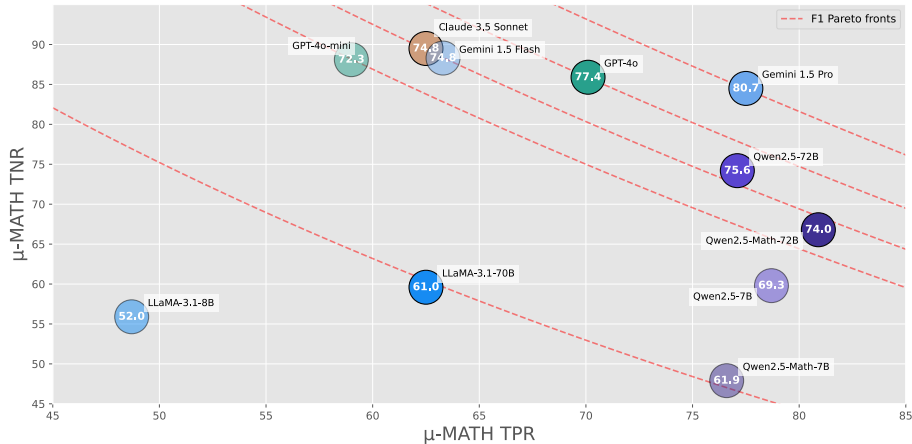


Figure 15: True Positive Rate vs True Negative Rate of judges on μ -MATH. The value inside of the marker denotes the macro F1-score.

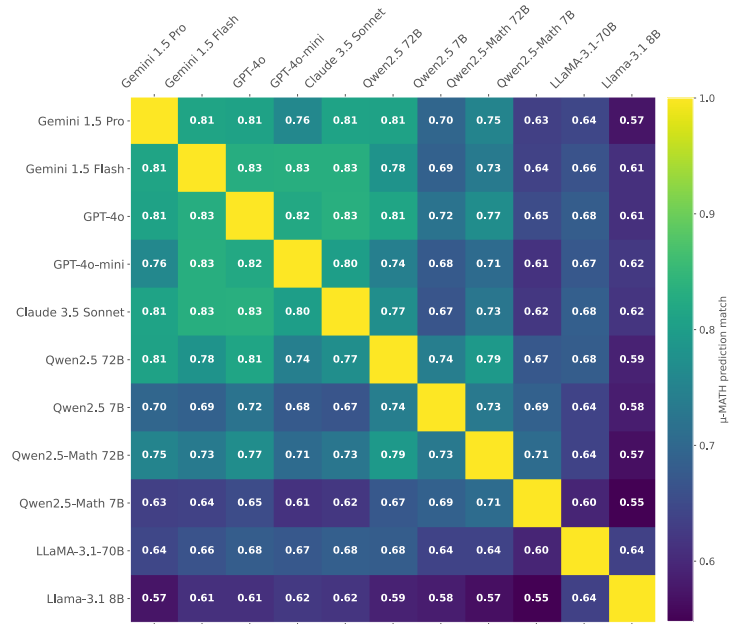


Figure 16: Agreement between different judges on μ -MATH as measured by predicted label coincidence ratio.

J μ -MATH JUDGMENT EXAMPLES

Example: Differential calculus.

U-MATH Problem:

Compute the derivative of the function

$$y = \sqrt{\frac{x^5 \cdot (2 \cdot x^6 + 3)}{\sqrt[3]{1 - 2 \cdot x}}}$$

Golden answer:

$$y' = \frac{-128 \cdot x^7 + 66 \cdot x^6 - 84 \cdot x + 45}{-24 \cdot x^8 + 12 \cdot x^7 - 36 \cdot x^2 + 18 \cdot x} \cdot \sqrt{\frac{x^5 \cdot (2 \cdot x^6 + 3)}{\sqrt[3]{1 - 2 \cdot x}}}$$

LLM-generated answer:

$$\sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1 - 2x}}} \cdot \frac{1}{2} \left(\frac{5}{x} + \frac{12x^5}{2x^6 + 3} + \frac{2}{3(1 - 2x)} \right)$$

Golden judge verdict: correct

Claude 3.5 Sonnet Judgment.

1. The solution's answer is: $\sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1 - 2x}}} \cdot \frac{1}{2} \left(\frac{5}{x} + \frac{12x^5}{2x^6 + 3} + \frac{2}{3(1 - 2x)} \right)$
2. The correct answer is: $\frac{-128 \cdot x^7 + 66 \cdot x^6 - 84 \cdot x + 45}{-24 \cdot x^8 + 12 \cdot x^7 - 36 \cdot x^2 + 18 \cdot x} \cdot \sqrt{\frac{x^5 \cdot (2 \cdot x^6 + 3)}{\sqrt[3]{1 - 2 \cdot x}}}$
3. To compare these expressions:
 - Both have the same factor $\sqrt{\frac{x^5 \cdot (2 \cdot x^6 + 3)}{\sqrt[3]{1 - 2 \cdot x}}}$
 - However, the other factors are completely different:
 - Solution has: $\frac{1}{2} \left(\frac{5}{x} + \frac{12x^5}{2x^6 + 3} + \frac{2}{3(1 - 2x)} \right)$
 - Correct answer has: $\frac{-128 \cdot x^7 + 66 \cdot x^6 - 84 \cdot x + 45}{-24 \cdot x^8 + 12 \cdot x^7 - 36 \cdot x^2 + 18 \cdot x}$
 - These expressions are not algebraically equivalent, as they have different degrees in numerator and denominator
4. No

1. The solution provides the derivative as:

$$\frac{dy}{dx} = \sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1-2x}}} \cdot \frac{1}{2} \left(\frac{5}{x} + \frac{12x^5}{2x^6 + 3} + \frac{2}{3(1-2x)} \right)$$

2. The provided correct answer is:

$$y' = \frac{-128x^7 + 66x^6 - 84x + 45}{-24x^8 + 12x^7 - 36x^2 + 18x} \cdot \sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1-2x}}}$$

We need to simplify the expression inside the parentheses of the solution and see if it matches the fraction in the correct answer.

First, let's rewrite the expression inside the parentheses:

$$\frac{5}{x} + \frac{12x^5}{2x^6 + 3} + \frac{2}{3(1-2x)}$$

To combine these fractions, we need a common denominator. The common denominator is

$$x(2x^6 + 3)(3(1-2x)) = 3x(2x^6 + 3)(1-2x)$$

Rewriting each term with this common denominator:

$$\begin{aligned} \frac{5}{x} &= \frac{5 \cdot 3(2x^6 + 3)(1-2x)}{3x(2x^6 + 3)(1-2x)} = \frac{15(2x^6 + 3)(1-2x)}{3x(2x^6 + 3)(1-2x)} \\ \frac{12x^5}{2x^6 + 3} &= \frac{12x^5 \cdot 3x(1-2x)}{3x(2x^6 + 3)(1-2x)} = \frac{36x^6(1-2x)}{3x(2x^6 + 3)(1-2x)} \\ \frac{2}{3(1-2x)} &= \frac{2 \cdot x(2x^6 + 3)}{3x(2x^6 + 3)(1-2x)} = \frac{2x(2x^6 + 3)}{3x(2x^6 + 3)(1-2x)} \end{aligned}$$

Combining these fractions:

$$\frac{15(2x^6 + 3)(1-2x) + 36x^6(1-2x) + 2x(2x^6 + 3)}{3x(2x^6 + 3)(1-2x)}$$

Expanding the numerator:

$$\begin{aligned} 15(2x^6 + 3)(1-2x) &= 30x^6 - 60x^7 + 45 - 90x \\ 36x^6(1-2x) &= 36x^6 - 72x^7 \\ 2x(2x^6 + 3) &= 4x^7 + 6x \end{aligned}$$

Adding these together:

$$30x^6 - 60x^7 + 45 - 90x + 36x^6 - 72x^7 + 4x^7 + 6x = -128x^7 + 66x^6 - 84x + 45$$

So the expression inside the parentheses simplifies to:

$$\frac{-128x^7 + 66x^6 - 84x + 45}{3x(2x^6 + 3)(1-2x)}$$

Therefore, the derivative is:

$$\sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1-2x}}} \cdot \frac{1}{2} \cdot \frac{-128x^7 + 66x^6 - 84x + 45}{3x(2x^6 + 3)(1-2x)} = \frac{-128x^7 + 66x^6 - 84x + 45}{6x(2x^6 + 3)(1-2x)} \cdot \sqrt{\frac{x^5 \cdot (2x^6 + 3)}{\sqrt[3]{1-2x}}}$$

Notice that:

$$-24x^8 + 12x^7 - 36x^2 + 18x = -6x(4x^7 - 2x^6 + 6x - 3) = -6x(2x^6 + 3)(2x - 1) = 6x(2x^6 + 3)(1 - 2x)$$

Thus, the two expressions are identical. Therefore, the solution is correct. The final answer is:

Yes

K FULL μ -MATH RESULTS

Model	μ -MATH					μ -MATH _{Qwen}					μ -MATH _{Llama}					μ -MATH _{GPT}					μ -MATH _{Gemini}				
	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV
Llama-3.1 8B	52.0	48.7	55.9	56.0	48.5	48.7	45.2	53.4	56.5	42.2	49.2	45.4	54.0	35.5	63.9	51.2	46.4	56.5	53.3	49.7	55.5	55.0	62.2	77.0	37.5
Qwen2.5 7B	69.3	78.7	59.8	69.3	70.8	62.4	75.5	49.1	66.5	60.0	72.3	78.4	70.1	59.4	85.3	68.3	81.4	55.7	66.3	73.7	69.1	79.4	59.8	82.0	55.7
Qwen2.5-Math 7B	61.9	76.6	47.9	62.9	63.9	59.7	71.0	48.3	64.7	55.4	63.8	85.6	51.7	49.7	86.5	57.2	75.0	41.2	57.7	60.7	63.8	77.8	50.0	78.2	49.4
GPT-4o-mini	72.3	59.0	88.1	85.1	65.1	69.3	56.1	87.1	85.3	59.8	76.2	59.8	90.2	77.3	80.1	70.4	56.4	86.3	81.4	64.9	69.6	63.0	87.8	92.2	50.7
Gemini 1.5 Flash	74.8	63.3	88.3	86.2	67.6	71.2	60.6	85.3	84.7	61.9	80.6	67.0	92.0	82.3	83.3	70.1	57.9	84.0	79.4	65.1	73.9	67.7	91.5	94.8	55.1
Llama-3.1-70B	61.0	62.5	59.6	64.1	57.9	56.0	58.7	53.4	62.8	49.2	57.0	63.9	54.0	43.7	72.9	69.4	67.1	71.8	71.8	67.1	58.8	61.4	61.0	78.4	40.7
Qwen2.5 72B	75.6	77.1	74.2	77.5	73.7	70.5	76.1	64.7	74.2	67.0	79.3	75.3	83.9	72.3	85.9	73.7	76.4	71.0	73.8	73.8	74.2	79.4	72.0	86.7	60.2
Qwen2.5-Math 72B	74.0	80.9	66.8	73.8	75.2	69.3	81.3	56.9	71.6	69.5	77.3	81.4	76.4	65.8	88.1	68.2	77.9	58.8	66.9	71.3	76.8	82.5	73.2	87.6	64.5
GPT-4o	77.4	70.1	85.9	85.1	71.3	74.2	67.1	83.6	84.6	65.5	81.8	71.1	90.8	81.2	84.9	77.5	72.1	83.2	82.1	73.6	72.6	70.4	82.9	90.5	54.8
Claude 3.5 Sonnet	74.8	62.5	89.5	87.3	67.4	70.8	58.1	87.9	86.5	61.1	77.9	59.8	93.1	82.9	80.6	72.2	57.9	88.5	84.4	66.3	73.8	70.9	85.4	91.8	56.0
Gemini 1.5 Pro	80.7	77.5	84.5	85.2	76.4	77.7	75.5	81.0	84.2	71.2	83.6	75.3	90.8	82.0	86.8	78.2	76.4	80.2	80.5	76.1	79.5	81.0	82.9	91.6	65.4

Table 8: Comparison of model’s ability to judge on μ -MATH benchmark, **with CoT prompting**; Macro F1-score (F1), True Positive Rate (TPR), True Negative Rate (TNR), Positive Predictive Value (PPV), and Negative Predictive Value (NPV), with F1 as the primary one are presented. Columns under μ -MATH represent the integral score over the entire benchmark, while μ -MATH_{<model>} columns denote subsets with solutions generated by a specific author model. **Bold** indicates the best result in each group.

Model	μ -MATH					μ -MATH _{Qwen}					μ -MATH _{Llama}					μ -MATH _{GPT}					μ -MATH _{Gemini}				
	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV	F1	TPR	TNR	PPV	NPV
Llama-3.1 8B	53.1	36.3	75.0	62.6	50.5	49.6	34.8	70.7	61.4	44.8	51.2	29.9	73.0	38.2	65.1	57.6	41.4	77.1	65.9	55.2	50.5	37.0	81.7	82.4	36.0
Qwen2.5 7B	67.0	65.9	68.4	70.7	63.5	60.5	61.3	60.3	67.4	53.8	72.4	71.1	75.3	61.6	82.4	66.4	67.1	65.6	67.6	65.2	65.0	66.1	69.5	83.3	47.1
Qwen2.5-Math 7B	61.2	72.8	49.9	62.7	61.4	56.7	67.7	45.7	62.5	51.5	64.0	80.4	55.2	50.0	83.5	58.5	68.6	48.9	58.9	59.3	61.2	76.2	46.3	76.6	45.8
GPT-4o-mini	69.2	50.8	92.0	88.1	61.8	61.7	43.2	87.9	82.7	53.7	78.5	58.8	94.8	86.4	80.5	69.8	50.7	92.4	87.7	63.7	64.3	52.9	91.5	93.5	45.7
Gemini 1.5 Flash	65.3	42.0	95.8	92.1	58.9	61.9	39.4	94.8	91.0	53.9	70.8	41.2	97.7	90.9	74.9	65.3	40.7	96.2	91.9	60.3	59.7	45.5	92.7	93.5	42.5
Llama-3.1-70B	68.2	60.9	76.7	75.2	63.0	63.8	58.7	70.7	72.8	56.2	70.2	60.8	79.3	62.1	78.4	69.8	58.6	82.4	78.1	65.1	64.4	64.6	70.7	83.6	46.4
Qwen2.5 72B	75.1	73.0	77.5	79.0	71.3	68.9	64.5	75.0	77.5	61.3	80.1	76.3	84.5	73.3	86.5	73.4	74.3	72.5	74.3	72.5	73.8	77.2	74.4	87.4	58.7
Qwen2.5-Math 72B	75.5	77.1	74.0	77.4	73.7	68.8	72.3	65.5	73.7	63.9	79.8	77.3	83.3	72.1	86.8	69.2	73.6	64.9	69.1	69.7	80.4	83.6	80.5	90.8	68.0
GPT-4o	74.2	61.4	89.3	86.9	66.7	68.2	56.8	83.6	82.2	59.1	78.9	61.9	93.1	83.3	81.4	75.8	65.0	87.8	85.0	70.1	70.5	62.4	91.5	94.4	51.4
Claude 3.5 Sonnet	68.1	46.8	95.2	91.9	60.8	64.1	43.2	94.0	90.5	55.3	71.8	41.2	99.4	97.6	75.2	68.1	45.7	95.4	91.4	62.2	63.4	53.4	87.8	91.0	45.0
Gemini 1.5 Pro	69.1	49.4	93.8	90.3	61.6	64.6	44.5	93.1	89.6	55.7	74.8	51.5	94.8	84.7	77.8	68.6	46.4	95.4	91.5	62.5	64.9	54.5	90.2	92.8	46.2

Table 9: Comparison of model’s ability to judge on μ -MATH benchmark, **with AutoCoT prompting**; Macro F1-score (F1), True Positive Rate (TPR), True Negative Rate (TNR), Positive Predictive Value (PPV), and Negative Predictive Value (NPV), with F1 as the primary one are presented. Columns under μ -MATH represent the integral score over the entire benchmark, while μ -MATH_{<model>} columns denote subsets with solutions generated by a specific author model. **Bold** indicates the best result in each group.